



# Scientific Software Design

Anshu Dubey

Argonne National Laboratory

Software Productivity Track, ATPESC 2020



See slide 2 for  
license details

# License, Citation and Acknowledgements



## License and Citation

- This work is licensed under a [Creative Commons Attribution 4.0 International License](#) (CC BY 4.0).
- The requested citation for the overall tutorial is: David E. Bernholdt, Anshu Dubey, Mark C. Miller, Katherine M. Riley, and James M. Willenbring, Software Productivity Track, in Argonne Training Program for Extreme Scale Computing (ATPESC), August 2020, online. DOI: [10.6084/m9.figshare.12719834](https://doi.org/10.6084/m9.figshare.12719834)
- Individual modules may be cited as *Speaker, Module Title*, in Software Productivity Track...

## Acknowledgements

- Additional contributors include: Patricia Grubel, Rinku Gupta, Mike Heroux, Alicia Klinvex, Jared O’Neal, David Rogers, Deborah Stevens
- This work was supported by the U.S. Department of Energy Office of Science, Office of Advanced Scientific Computing Research (ASCR), and by the Exascale Computing Project (17-SC-20-SC), a collaborative effort of the U.S. Department of Energy Office of Science and the National Nuclear Security Administration.
- This work was performed in part at the Argonne National Laboratory, which is managed by UChicago Argonne, LLC for the U.S. Department of Energy under Contract No. DE-AC02-06CH11357.
- This work was performed in part at the Oak Ridge National Laboratory, which is managed by UT-Battelle, LLC for the U.S. Department of Energy under Contract No. DE-AC05-00OR22725.
- This work was performed in part at the Lawrence Livermore National Laboratory, which is managed by Lawrence Livermore National Security, LLC for the U.S. Department of Energy under Contract No. DE-AC52-07NA27344.
- This work was performed in part at Sandia National Laboratories. Sandia National Laboratories is a multi-mission laboratory managed and operated by National Technology and Engineering Solutions of Sandia, LLC., a wholly owned subsidiary of Honeywell International, Inc., for the U.S. Department of Energy’s National Nuclear Security Administration under contract DE-NA0003525.

# Architecting scientific codes

## Desirable Characteristics and Why They are Challenging

Extensibility

Well defined structure and  
modules  
Encapsulation of  
functionalities

# Architecting scientific codes

## Desirable Characteristics and Why They are Challenging

Extensibility

Well defined structure and  
modules  
Encapsulation of  
functionalities

Same data layout not  
good for all solvers. Many  
corner cases. Necessary  
lateral interactions

# Architecting scientific codes

## Desirable Characteristics and Why They are Challenging

### Extensibility

Well defined structure and modules  
Encapsulation of functionalities

### Performance

Spatial and temporal locality of data  
Minimizing data movement  
Maximizing scalability

Same data layout not good for all solvers. Many corner cases. Necessary lateral interactions

# Architecting scientific codes

## Desirable Characteristics and Why They are Challenging

### Extensibility

Well defined structure and modules  
Encapsulation of functionalities

### Performance

Spatial and temporal locality of data  
Minimizing data movement  
Maximizing scalability

Same data layout not good for all solvers. Many corner cases. Necessary lateral interactions

Low arithmetic intensity solvers with hard dependencies. Proximity and work distribution at cross purposes

# Architecting scientific codes

## Desirable Characteristics and Why They are Challenging

Portability

General solutions that  
work without significant  
manual intervention  
across platforms

# Architecting scientific codes

## Desirable Characteristics and Why They are Challenging

Portability

General solutions that work without significant manual intervention across platforms

Tremendous platform heterogeneity  
A version for each class of device => combinatorial explosion

# Architecting scientific codes

## Desirable Characteristics and Why They are Challenging

Portability

General solutions that work without significant manual intervention across platforms

Verifiability and Maintainability

Clean code  
Documentation  
Comprehensive testing

Tremendous platform heterogeneity  
A version for each class of device => combinatorial explosion

# Architecting scientific codes

## Desirable Characteristics and Why They are Challenging

Portability

General solutions that work without significant manual intervention across platforms

Verifiability and Maintainability

Clean code  
Documentation  
Comprehensive testing

Tremendous platform heterogeneity  
A version for each class of device => combinatorial explosion

Wrong incentives  
Designing good tests is hard

# Architecting scientific codes

## Taming the Complexity: Separation of Concerns

**Subject of research**  
Model  
Numerics

**More Stable**  
Discretization  
I/O  
Parameters

# Architecting scientific codes

## Taming the Complexity: Separation of Concerns

**Subject of research**  
Model  
Numerics

Treat differently

**More Stable**  
Discretization  
I/O  
Parameters

# Architecting scientific codes

## Taming the Complexity: Separation of Concerns

**Subject of research**  
Model  
Numerics

**Client Code**  
Mathematically complex

Treat differently

**More Stable**  
Discretization  
I/O  
Parameters

**Infrastructure**  
Data structures  
and movement

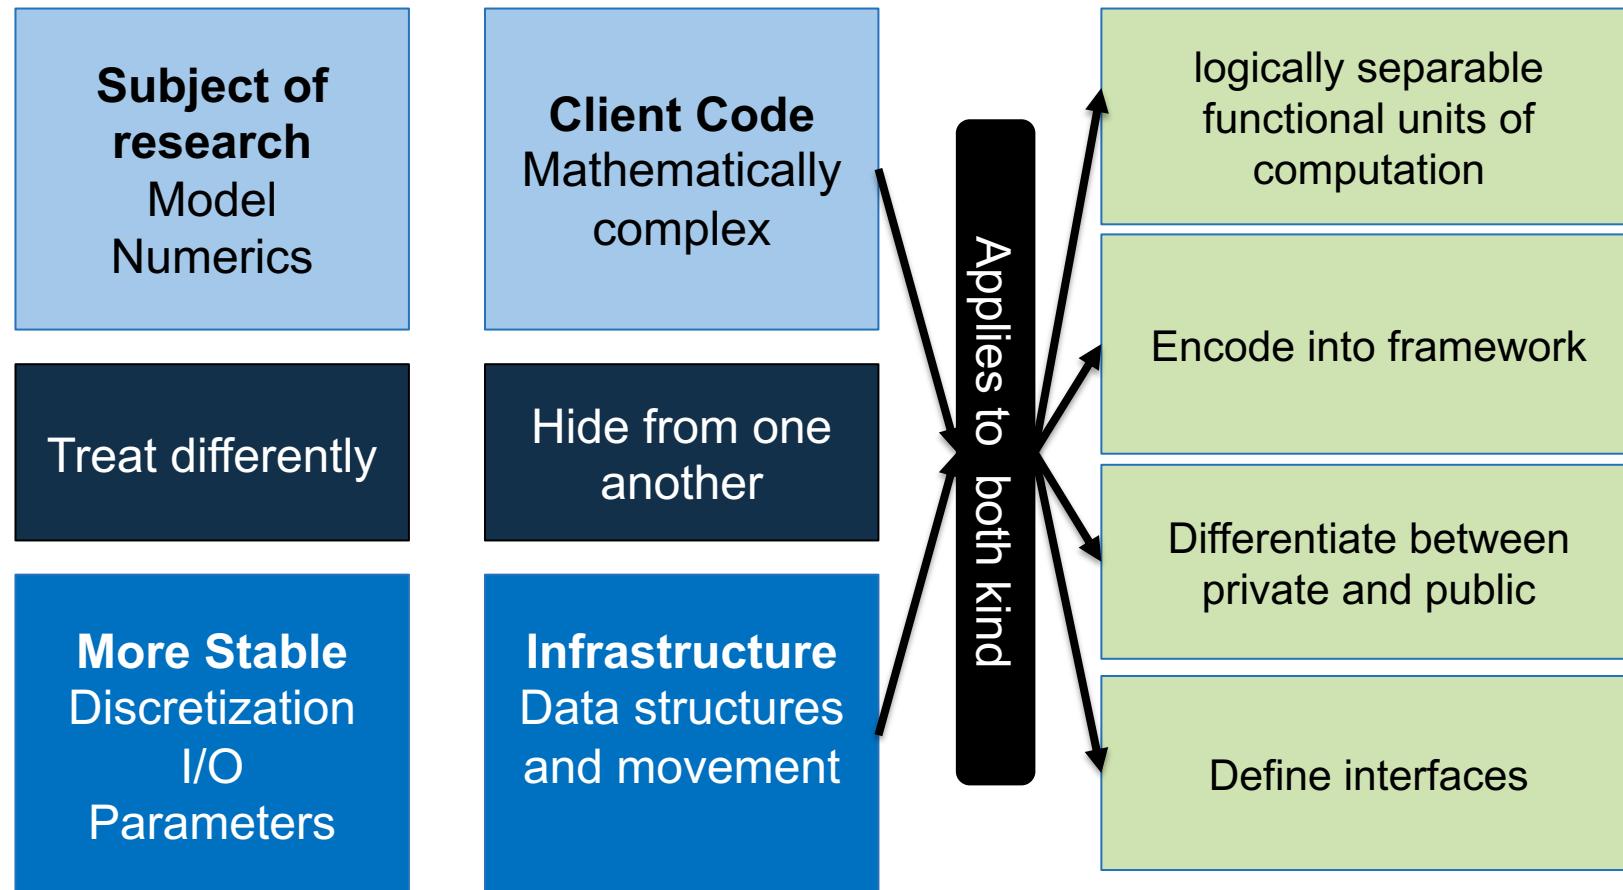
# Architecting scientific codes

## Taming the Complexity: Separation of Concerns

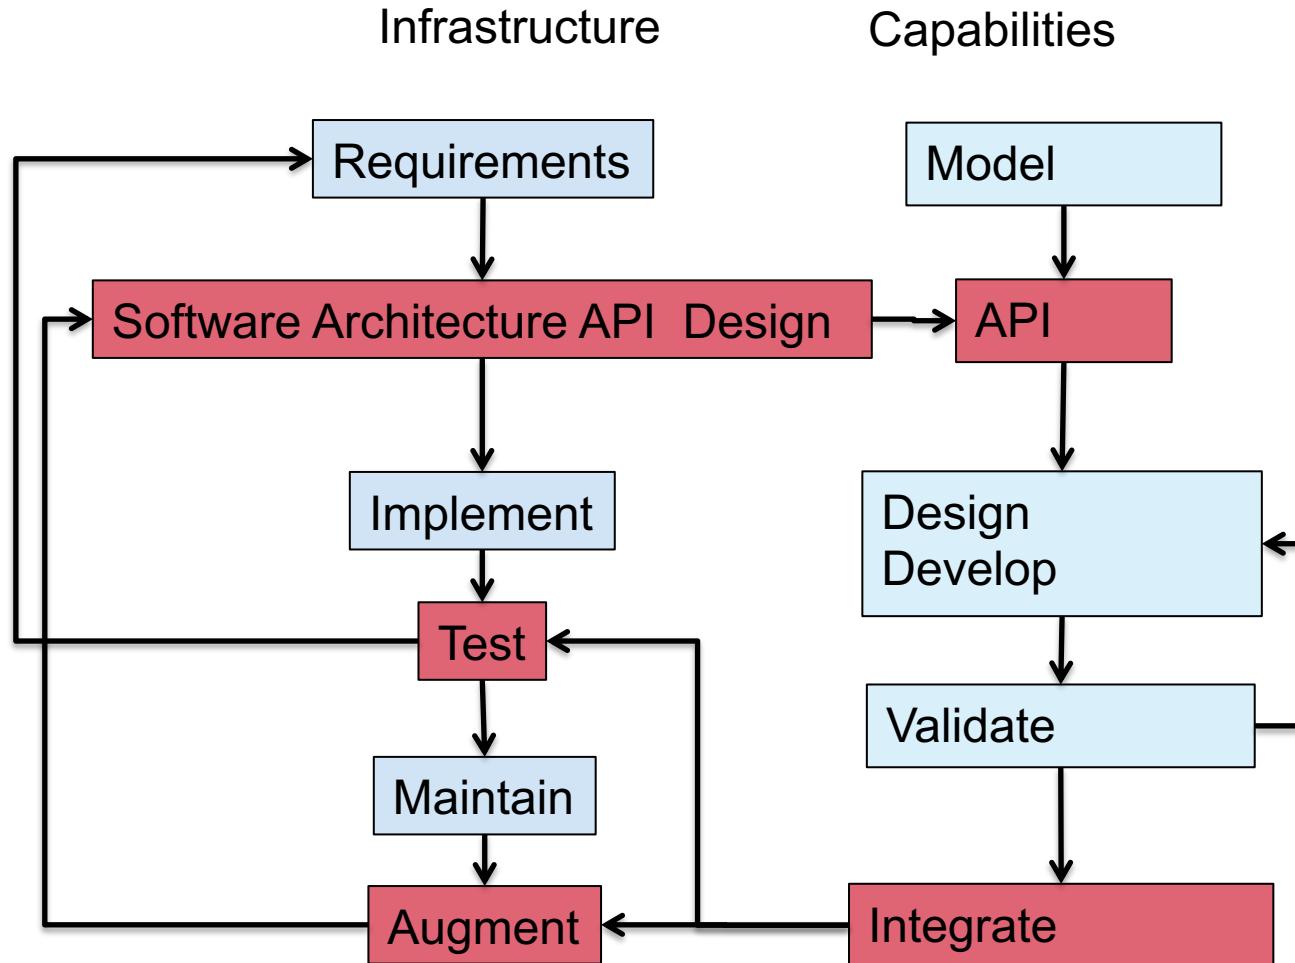
<b>Subject of research</b> Model Numerics	<b>Client Code</b> Mathematically complex
Treat differently	Hide from one another
<b>More Stable</b> Discretization I/O Parameters	<b>Infrastructure</b> Data structures and movement

# Architecting scientific codes

## Taming the Complexity: Separation of Concerns



# A Design Model for Separation of Concerns



# Design Considerations

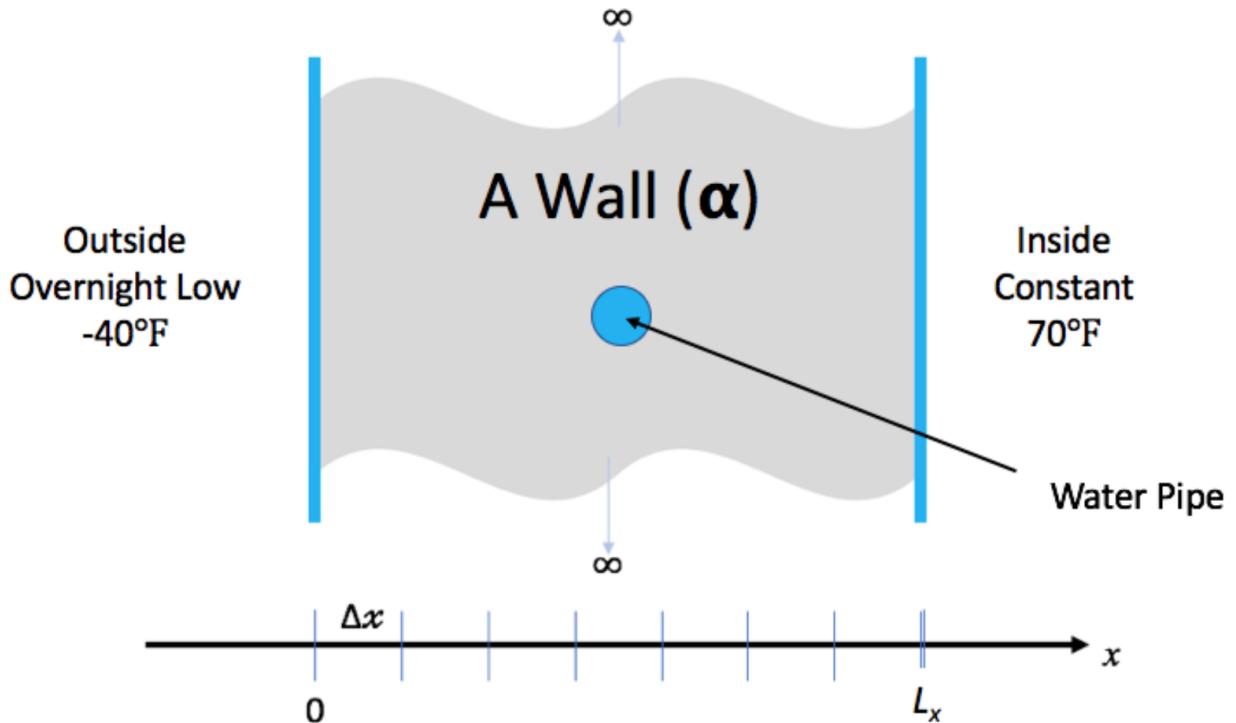
- Infrastructure design
  - Take time to discuss, iterate over requirements and specification
  - Keep end users involved
    - Not doing so leaves possible options on the table
- Simple is better
  - Flexibility Vs transparent to the user
    - Flexibility wins

# Design Considerations

- Infrastructure design
  - Take time to discuss, iterate over requirements and specification
  - Keep end users involved
    - Not doing so leaves possible options on the table
  - Keep API independent of numerics
- Simple is better
  - Flexibility Vs transparent to the user
    - Flexibility wins
- Model/numerics design
  - Abstract away the infrastructure knowledge as much as possible
  - Encapsulate
  - Let model needs guide API
  - Design flexible API to accommodate quick upgrades to methods
- Simple is better
  - Flexibility Vs transparent to the user
    - Flexibility wins

# The Running Example

Lets say you live in a house with exterior walls made of a single material of thickness,  $\$L_x\$$ . Inside the walls are some water pipes as pictured below.



[https://github.com/bett  
erscientificsoftware/hel  
lo-numerical-world-  
atpesc-2020](https://github.com/betterscientificsoftware/hello-numerical-world-atpesc-2020)

You keep the inside temperature of the house always at 70 degrees F. But, there is an overnight storm coming. The outside temperature is expected to drop to -40 degrees F for 15.5 hours. Will your pipes freeze before the storm is over?

# Problem Specification - Design Considerations

- Specification
  - Solve heat equation with some initial and boundary conditions
  - Apply different integration methods

- What is infrastructure here?
  - Discretization/ State
  - Verification
  - I/O
  - Application of initial conditions
  - Runtime parameters
  - Comparison

- What is model here?
  - Initial conditions
  - Boundary conditions
  - Integration

# Infrastructure API

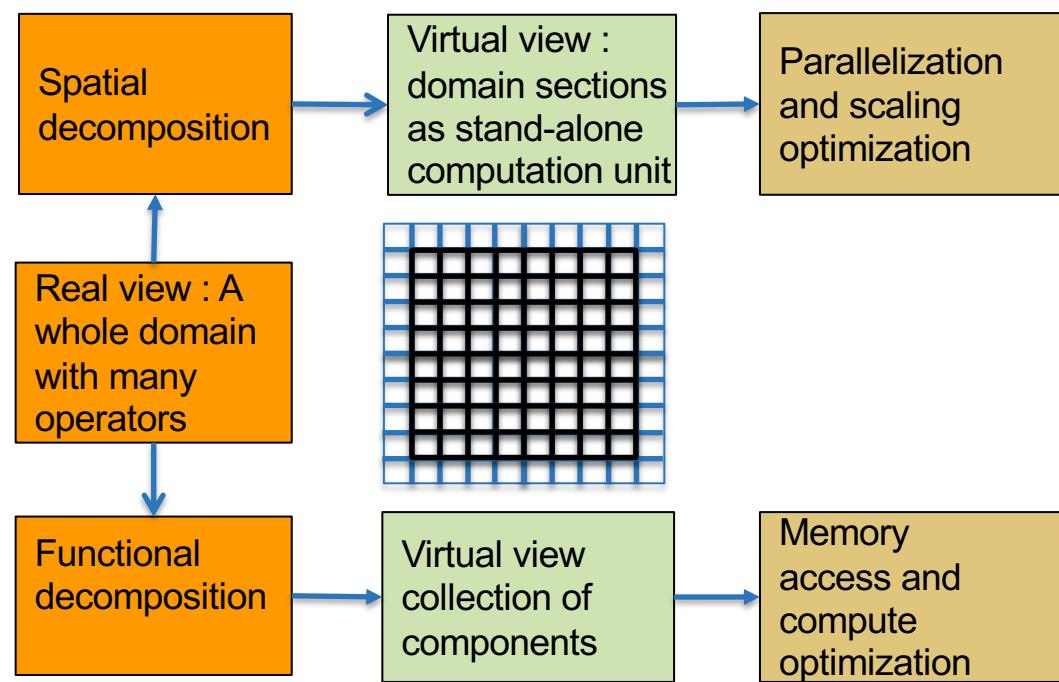
- `process_args(int argc, char **argv)`
- `static void initialize(void)`
- `void copy(int n, double *dst, double const *src)`
- `void write_array(int t, int n, double dx, double const *a)`
- `void set_initial_condition(int n, double *a, double dx, char const *ic)`

# Numerics API

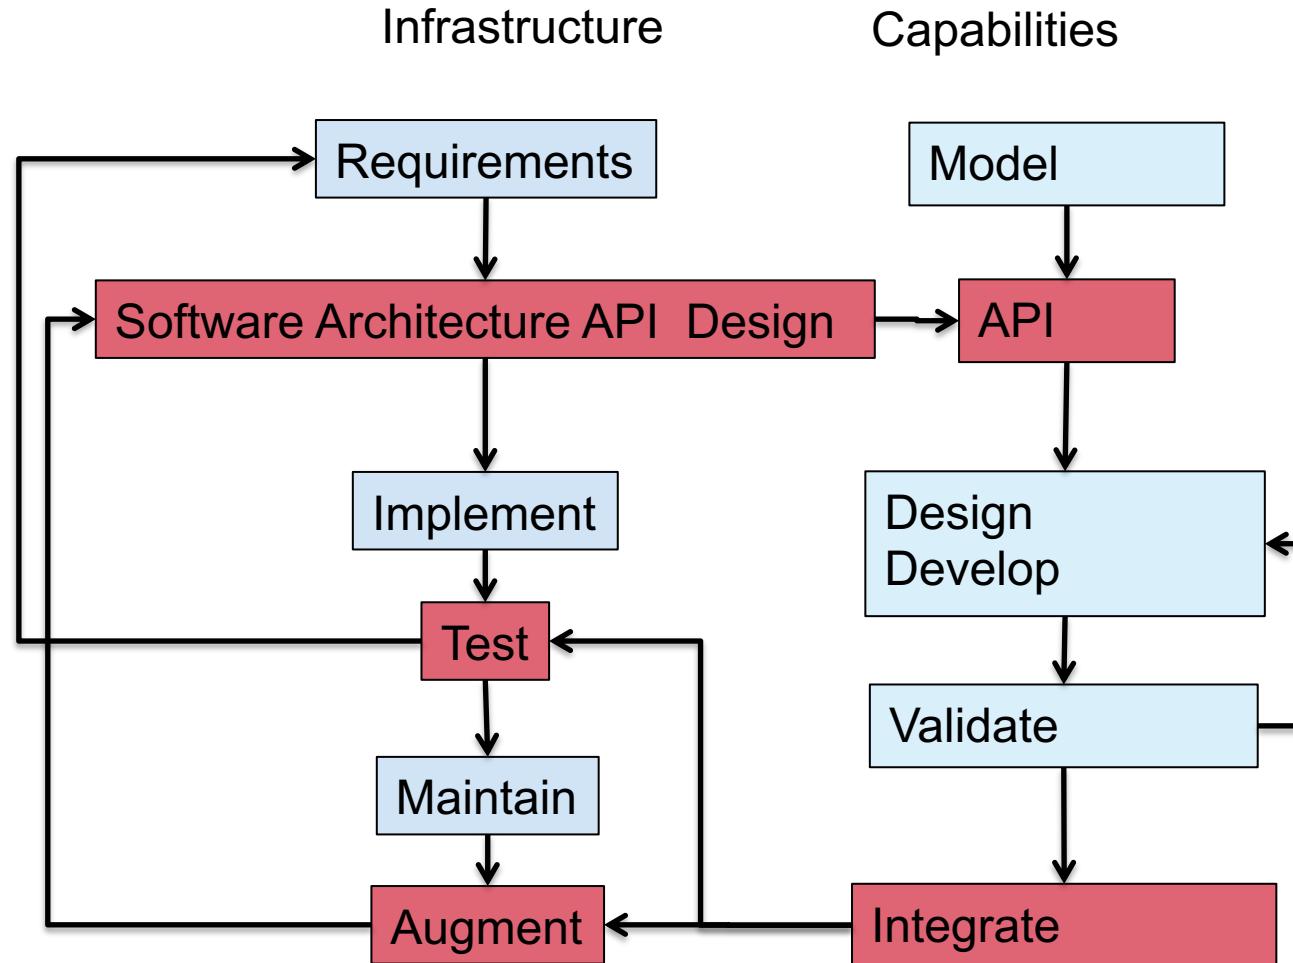
- double l2\_norm(int n, double const \*a, double const \*b)
- static void r83\_np\_fa(int n, double \*a)
- static void r83\_np\_sl ( int n, double const \*a\_lu, double const \*b, double \*x)
- bool update\_solution\_crankn(int n, double \*curr, double const \*last, double const \*cn\_Amat, double bc\_0, double bc\_1)
- bool update\_solution\_upwind15(int n, double \*curr, double const \*last, double alpha, double dx, double dt, double bc\_0, double bc\_1)
- void compute\_exact\_solution(int n, double \*a, double dx, char const \*ic, double alpha, double t, double bc0, double bc1)
- bool update\_solution\_ftcs( int n, double \*uk1, double const \*uk0, double alpha, double dx, double dt, double bc0, double bc1)

# Example: Architecting Multiphysics PDEs

- Virtual view of functionalities
- Decomposition into units and definition of interfaces



# A Design Model for Separation of Concerns

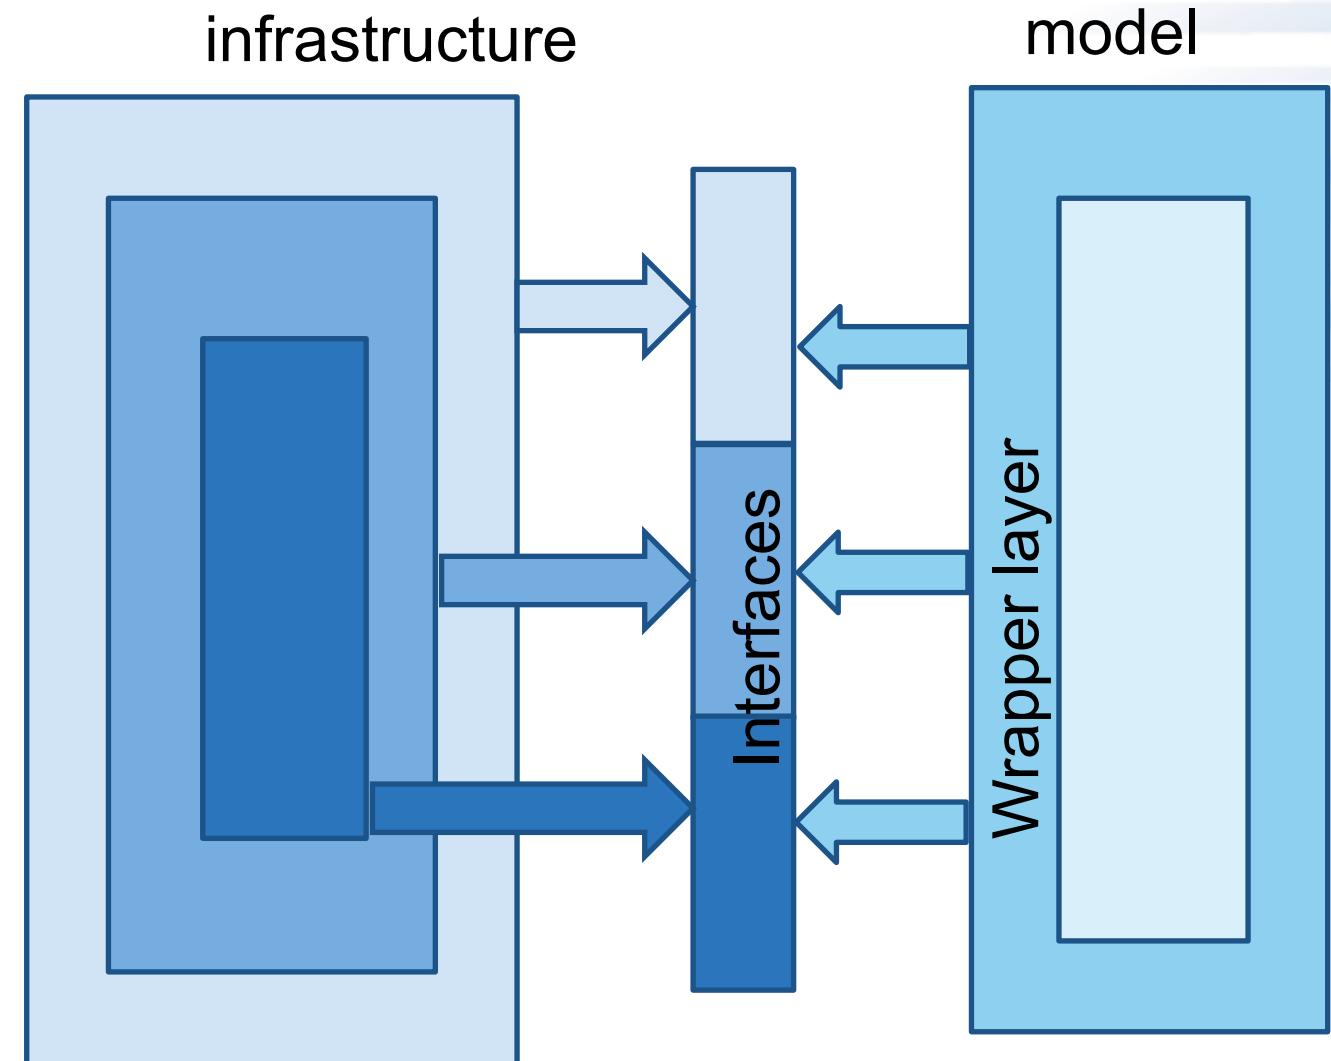


This worked with distributed memory parallelization model

No longer sufficient  
needs refinement

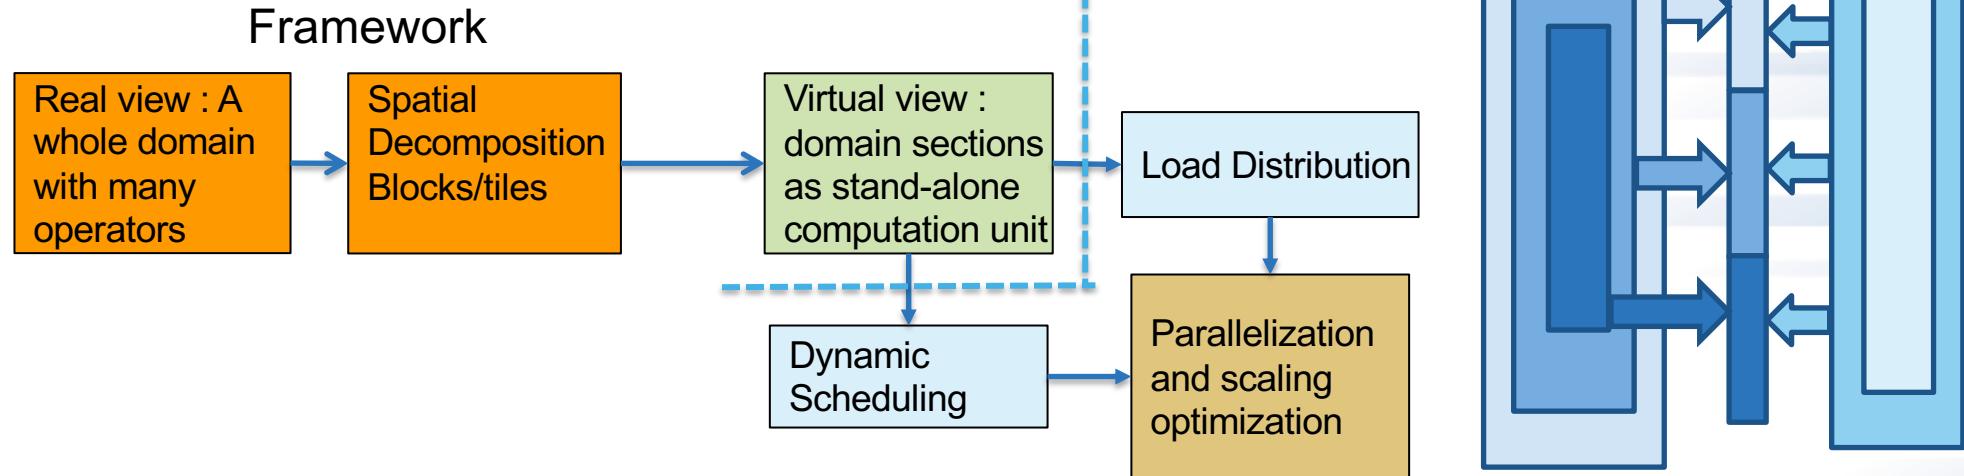
# Additional Considerations for Infrastructure

- Configurability
  - Components or kernels
  - Levels of access (hierarchical)
  - Layered API
- Task orchestration
  - Mapping tasks to devices
  - CPU, accelerators, specialized devices
  - Managing data movement between devices



# Example: Architecting Multiphysics PDEs

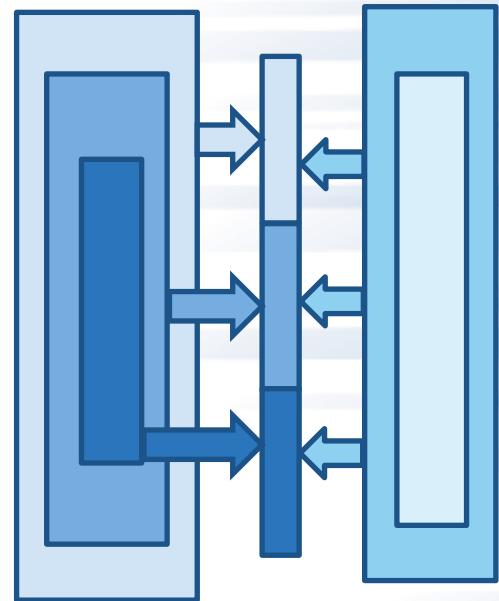
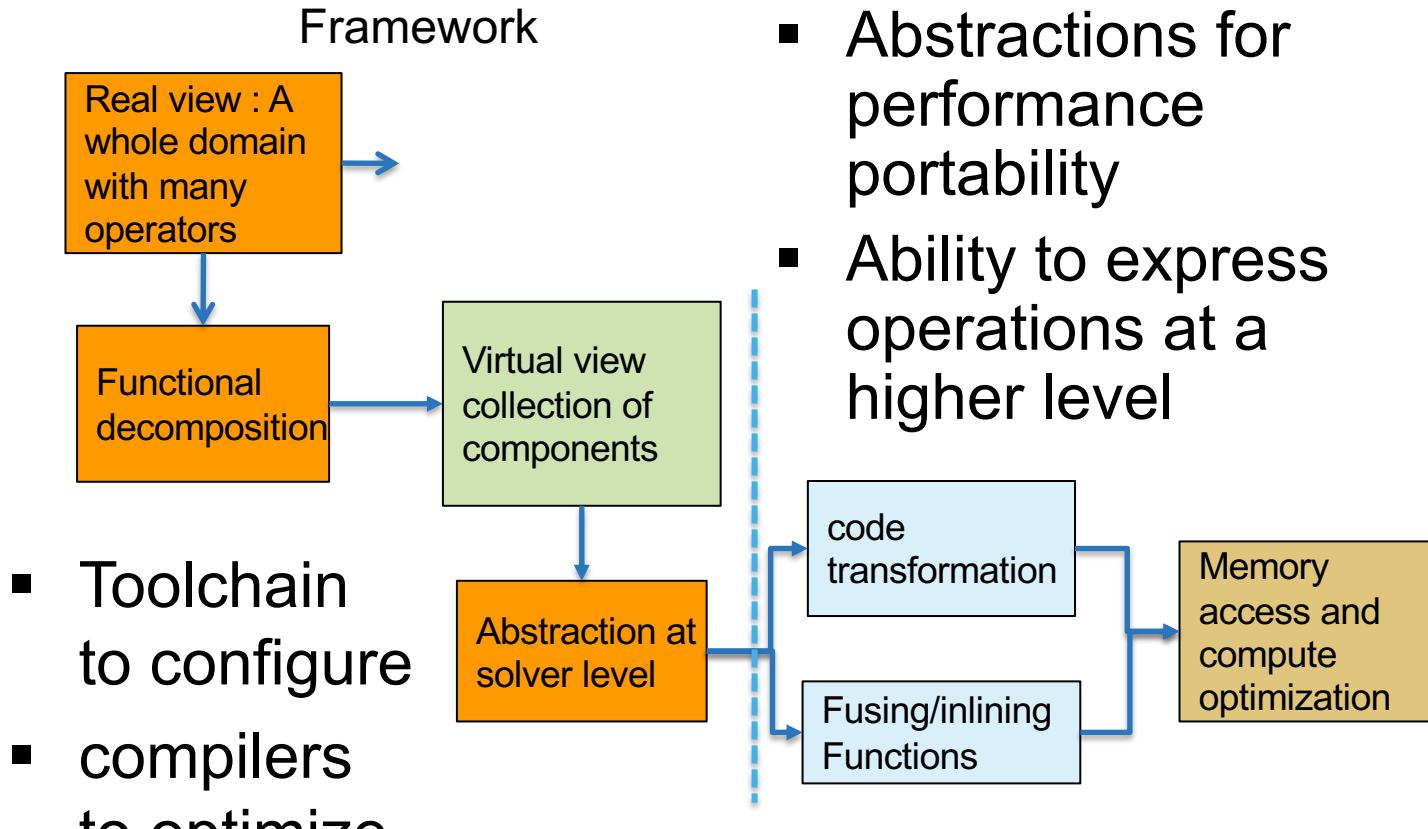
## Separation of Concerns, Tasks



- load balancing, work redistribution
- Meta-information about domain sections
- Possible asynchronous synchronization at block and operator level
- No compute optimization here

# Example: Architecting Multiphysics PDEs

## composition



# Other Considerations

- Leverage existing software
  - Libraries may have better solvers
    - Off-load expertise and maintenance
  - Examine the interoperability constraints
    - Many times the cost is justified even if there is more data movement
- More available packages are attempting to achieve interoperability
  - See if a combination meets your requirements
- May be worthwhile to let the library dictate data layout if the corresponding operations dominate

Institute a rigorous verification regime at the outset

# TAKEAWAYS

- DIFFERENTIATE BETWEEN SLOW CHANGING AND FAST CHANGING COMPONENTS OF YOUR CODE
  - TAKE YOUR TIME TO UNDERSTAND THE REQUIREMENTS OF YOUR INFRASTRUCTURE
  - IMPLEMENT SEPARATION OF CONCERNS
  - DESIGN WITH PORTABILITY, EXTENSIBILITY, REPRODUCIBILITY AND MAINTAINABILITY IN MIND
  - LEVERAGE EXISTING CAPABILITIES WHERE POSSIBLE
- .....QUESTIONS ?

# Agenda

Time (Central TZ)	Module	Topic	Speaker
9:30am-9:45am	00	Introduction	David E. Bernholdt, ORNL
9:45am-10:15am	01	Overview of Best Practices in HPC Software Development	Katherine M. Riley, ANL
10:15am-10:45am	02	Agile Methodologies	James M. Willenbring, SNL
10:45am-11:00am	03	Git Workflows	James M. Willenbring, SNL
<i>11:00am-11:15am</i>		<i>Break (and Q&amp;A with speakers)</i>	
11:15am-12:00pm	04	Software Design	Anshu Dubey, ANL
12:00pm-12:45pm	05	Software Testing	Anshu Dubey, ANL
<i>12:45pm-1:45pm</i>		<i>Lunch (and Q&amp;A with speakers)</i>	
1:45pm-2:00pm	06	Agile Methodologies Redux	James M. Willenbring, SNL
2:00pm-3:00pm	07	Refactoring	Anshu Dubey, ANL
<i>3:00pm-3:15pm</i>		<i>Break (and Q&amp;A with speakers)</i>	
3:15pm-3:45pm	08	Continuous Integration	Mark C. Miller, LLNL
3:45pm-4:30pm	09	Reproducibility	David E. Bernholdt, ORNL
4:30pm-4:45pm	10	Summary	David E. Bernholdt, ORNL

