

# Reproducible Workflows

*Presented by*

**COLABS: Collaboration  
for Better Software for  
Science**

Gregory R. Watson (he/him)  
Oak Ridge National Laboratory

*In collaboration with*



Software Practices for Reproducible Science tutorial @ ACM-REP 2024

Contributors: Patricia A. Grubel (LANL), David M. Rogers (ORNL),  
Gregory R. Watson (ORNL)

*With prior support from*



See slide 2 for  
license details

# License, Citation and Acknowledgements

## License and Citation

- This work is licensed under a [Creative Commons Attribution 4.0 International License](https://creativecommons.org/licenses/by/4.0/) (CC BY 4.0).
- **The requested citation the overall tutorial is:** Anshu Dubey and Gregory R. Watson, Software Practices for Reproducible Science tutorial, in 2024 ACM Conference on Reproducibility and Replicability (ACM-REP), Rennes, France and online, 2024. DOI: [10.6084/m9.figshare.26019469](https://doi.org/10.6084/m9.figshare.26019469).
- Individual modules may be cited as *Speaker, Module Title*, in *Tutorial Title*, ...



## Acknowledgements

- This work was supported by the U.S. Department of Energy Office of Science, Office of Advanced Scientific Computing Research (ASCR), and by the Exascale Computing Project (17-SC-20-SC), a collaborative effort of the U.S. Department of Energy Office of Science and the National Nuclear Security Administration.
- This work was supported by the U.S. Department of Energy, Office of Science, Office of Advanced Scientific Computing Research, Next-Generation Scientific Software Technologies (NGSST) program.
- This work was performed in part at the Argonne National Laboratory, which is managed by UChicago Argonne, LLC for the U.S. Department of Energy under Contract No. DE-AC02-06CH11357.
- This work was performed in part at the Lawrence Livermore National Laboratory, which is managed by Lawrence Livermore National Security, LLC for the U.S. Department of Energy under Contract No. DE-AC52-07NA27344.
- This work was performed in part at the Los Alamos National Laboratory, which is managed by Triad National Security, LLC for the U.S. Department of Energy under Contract No. 89233218CNA000001.
- This work was performed in part at the Oak Ridge National Laboratory, which is managed by UT-Battelle, LLC for the U.S. Department of Energy under Contract No. DE-AC05-00OR22725.
- This work was performed in part at Sandia National Laboratories. Sandia National Laboratories is a multi-mission laboratory managed and operated by National Technology and Engineering Solutions of Sandia, LLC., a wholly owned subsidiary of Honeywell International, Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.

# Introduction

- A *workflow* can be any series of activities needed to complete a task
- In scientific computing, a workflow can help with reproducibility by
  - Describing in a formal way how to undertake activities for scientific problem-solving
  - Might be as simple as using a script, or as complex as using a domain specific workflow description language
- A *workflow system* can help with reproducibility by
  - Providing the infrastructure required to execute the activities described by a workflow
  - Tracking the workflow and data provenance
- Workflow concepts and tools make this job easier and will be discussed in this module


# Outline

- Workflow concepts
- Example workflow systems
- Common workflow patterns
- Choosing a workflow system
- Reproducibility issues and containers

# What are scientific workflows?

HPC workflow == one or more steps require HPC

- Scientific applications often consist of a series of computational or data manipulation steps
- Each step may be written in different languages
  - May use different programming models, numeric methods and libraries
- Steps often have complex dependencies between them
- Steps must be done in a specific sequence to produce correct output
  - Can include cycles
- Steps may produce/consume large amounts of data
  - Data may need to be available from one step to another
- Steps may require interaction from the user



# Reasons to use a workflow system

- I have a large complex problem that can be broken down into a number of distinct steps
- I need to scale to a large jobs/systems
- I want to document steps / run method
- I need to be able to track and restart failed components
- I want to be able to reproduce the results or have someone else reproduce them
- I need to track provenance of task execution and data movement
- My application requires building a complicated environment

# Workflow activities and components

- **Resource provisioning**

- Each step of the workflow needs to execute on compute resources
- These may be preconfigured (e.g. a supercomputer) or dynamically allocated (e.g. cloud)
- The execution environment needs to be established (e.g. libraries, runtimes, hardware drivers, etc.)

- **Data management**

- Steps may consume data – must be available before execution commences
- Steps may produce data – how much, where will it be stored?
- What to do with the data from the workflow?

- **Job submission**

- Access to HPC systems is typically through a batch system
- These require scripts that determine what resources are required for the job to run

- **Workflow management**

- Orchestrates the flow of activities needed to carry out the computation
- Can be driven by data or by tasks

- **Workflow description**

- Domain specific language describing workflow
- Many workflow systems employ Common Workflow Language (CWL)

# Example workflow systems

<https://s.apache.org/existing-workflow-systems>

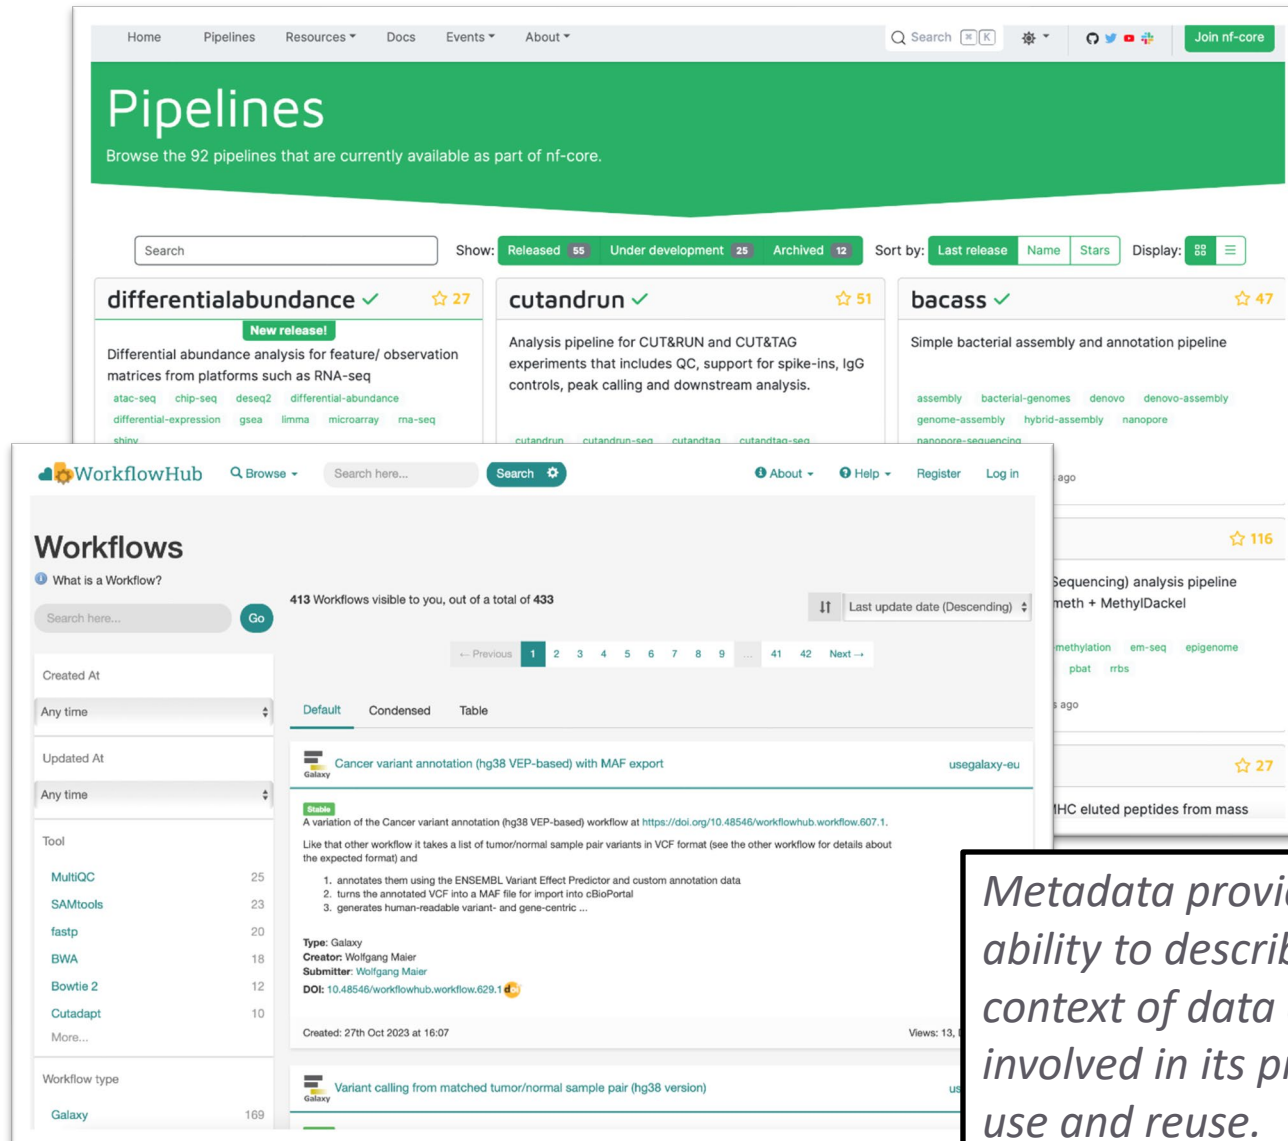
*Distributed*

*In situ*





# Workflow Repositories and Registries



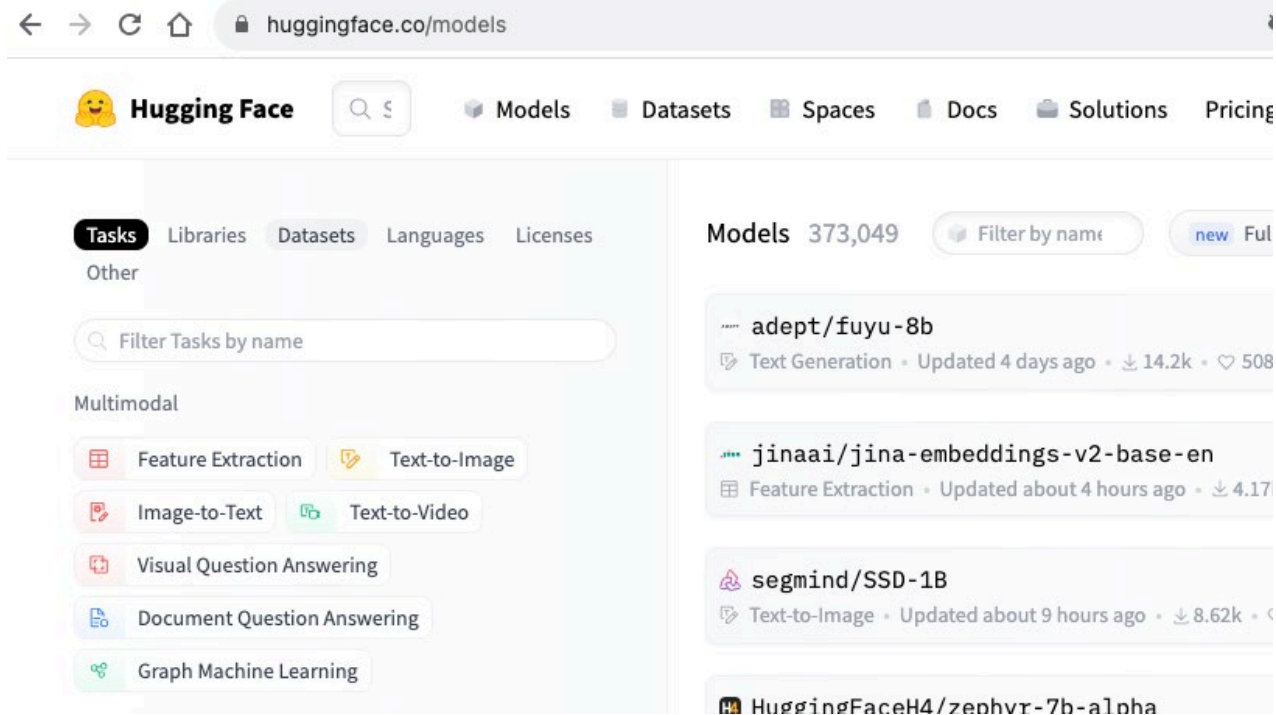
- Registry: indexes or catalogs of workflow metadata
- Repository: workflow descriptions and metadata are indexed and stored
- Metadata framework: help support reproducibility by managing metadata
- Data repository: permanently stores and manages data; provides digital object identifier (DOI)

Implementations:

<https://nf-co.re>  
<https://workflowhub.eu>  
<https://zenodo.org>

*Metadata provides the ability to describe the context of data and entities involved in its production, use and reuse.*


# Data / Code Combinator



- Publish data and structured code from within your code
- Offers tracking functionality for provenance (code / experiment progression)
- Popular for AI/ML models, where it facilitates model comparisons

Implementations:

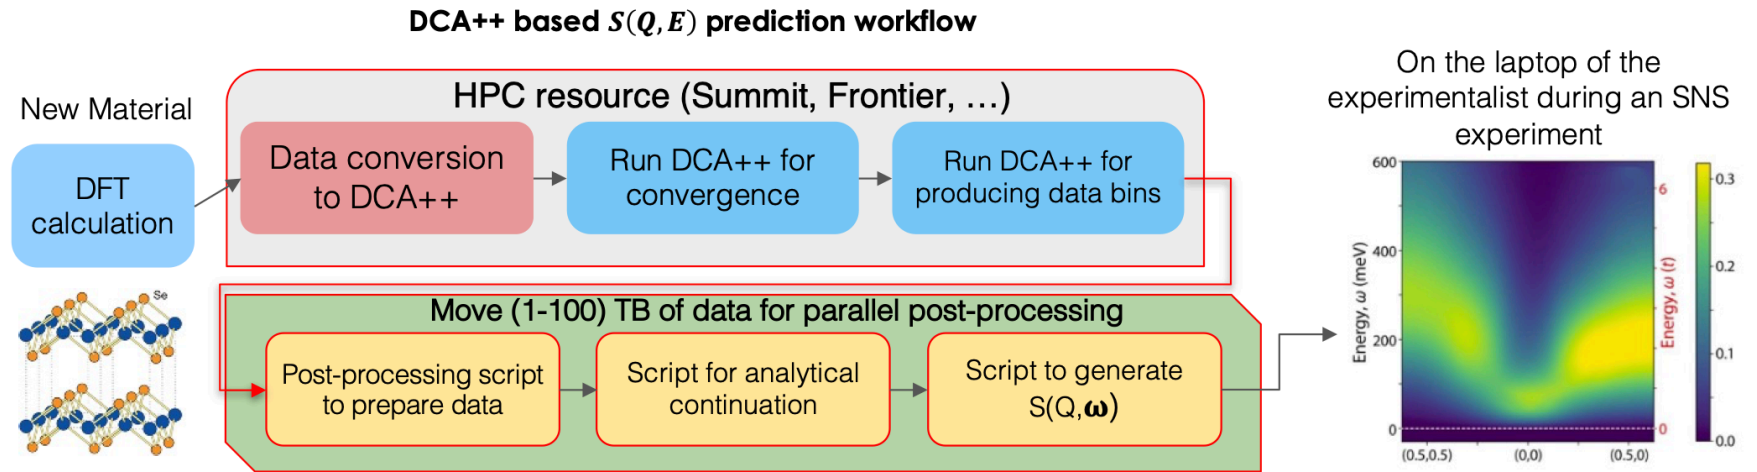
pypi / dvc ([DVC.org](https://dvc.org))  
Weights & Biases ([wandb.ai](https://wandb.ai))  
[huggingface.co](https://huggingface.co)



# Common scientific workflow patterns

## High-performance Simulation and Modeling

Figure 4. Diagram of the  $S(Q, \omega)$  prediction workflow. The workflow begins with a DFT calculation based on the new material of interest to build a low-energy effective model. A first DCA++ simulation is performed for convergence to self-consistency. For the problems of interest here, which require  $O(1000)$  bins of data for the four-point vertex function, a second DCA++ simulation is run to produce the binned data. The relevant observables are extracted and manipulated into a convenient form using several Python scripts. This includes the extraction of the vertex function  $\Gamma$ , the solution of the Bethe-Salpeter equation to generate  $S(Q, \omega)$  on the imaginary frequency axis, and the analytic continuation to perform the rotation to the real frequency axis. The final product of the Python analysis is the prediction for  $S(Q, \omega)$ .



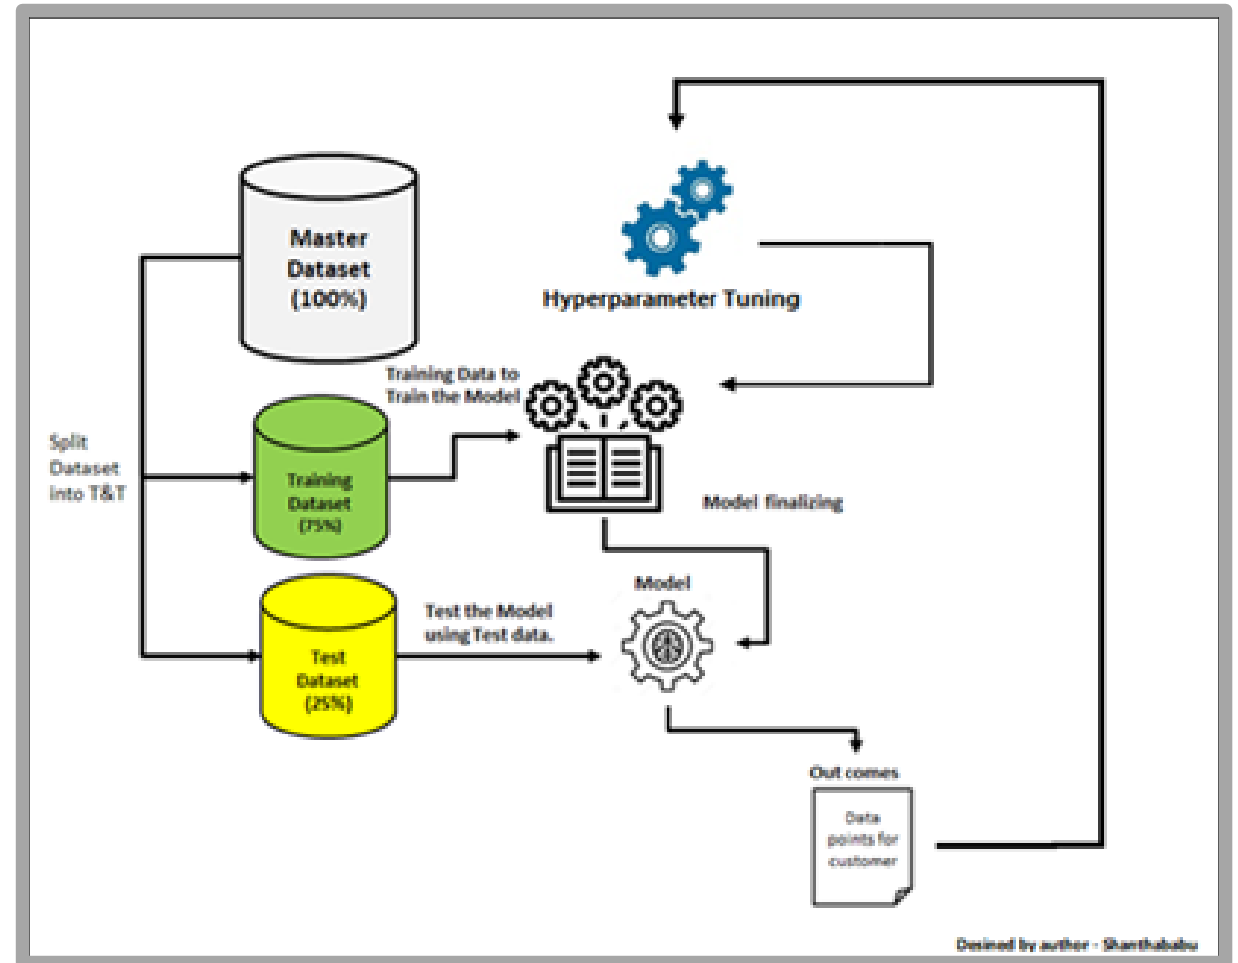
- Setup: mesh generation, pre-computation
- Simulate: physics calculations
- Down-sample: data reduction
- Post-process: format for visualization tools
- Visualize: inspect the data

Source: [https://doi.org/10.1007/978-3-031-23606-8\\_9](https://doi.org/10.1007/978-3-031-23606-8_9)

# Common scientific workflow patterns

## High-performance AI

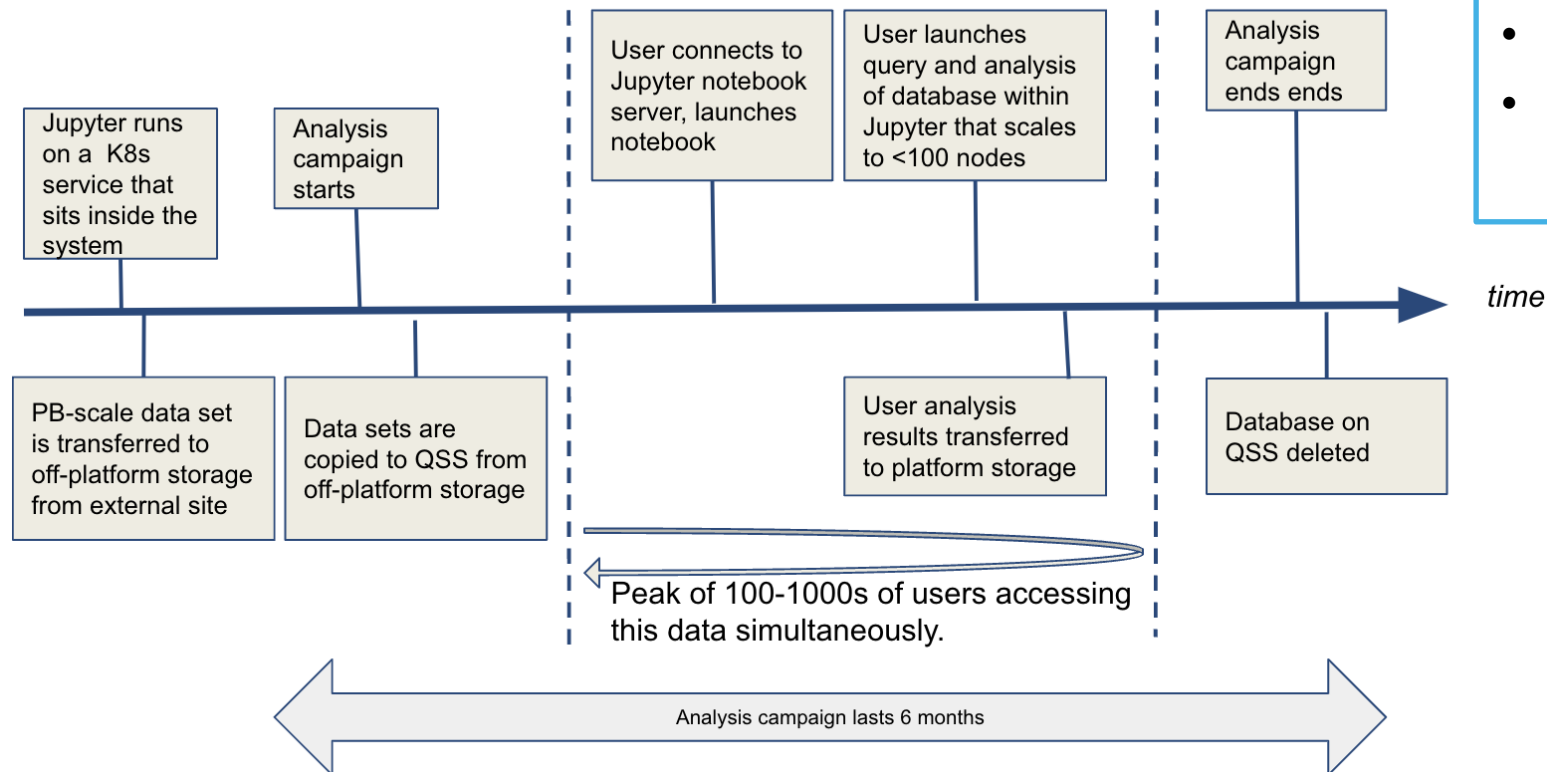
- Training: AI model training using scientific data
- Inference: Use the model to make predications
- Hyperparameter optimization: Find the best combination of hyperparameters for the problem



Source: <https://www.analyticsvidhya.com/blog/2022/02/a-comprehensive-guide-on-hyperparameter-tuning-and-its-techniques/>

# Common scientific workflow patterns

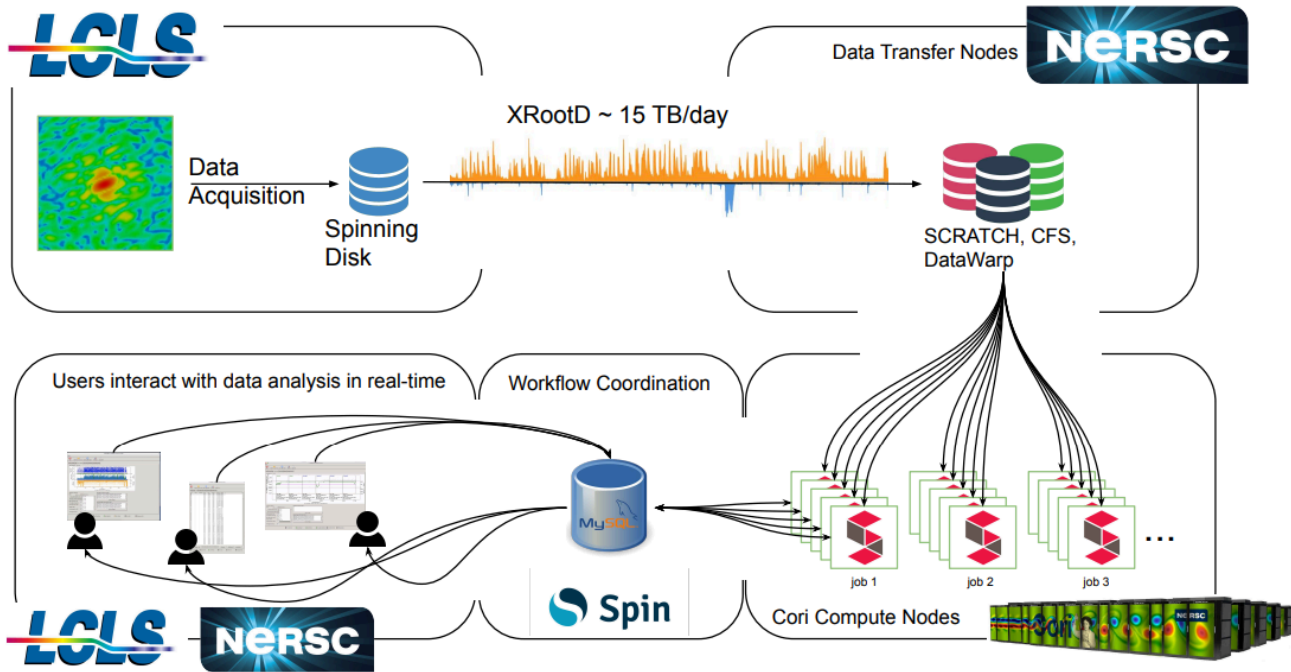
## Scientific Data Lifecycle



Source: <https://www.nersc.gov/assets/NERSC-10/Workflows-Archetypes-White-Paper-v1.0.pdf>

# Common scientific workflow patterns

## Real-Time / Streaming



- Data transport
- Compression
- Visualization
- Buffering / final storage

Implementations:

XRootD  
pypi / ray  
Apache kafka  
Globus  
custom...

Fig. 1. Sketch of the super-facility workflow: *Top*: Data are automatically transferred from the LCLS spinning-disk storage system via XRootD to NERSC's Scratch file system (the orange and blue spikes show the data transfer rate – approx. 1.3-2.6 GB/s – over the ESNet network, with each spike being a completed run). *Bottom*: On NERSC the CCTBX workers (running in Shifter containers on the Cori compute nodes) automatically analyze new data on Scratch, using the DataWarp burst buffer as a cache. Users at LCLS and NERSC connect to a MySQL database hosted at NERSC to orchestrate the workers, review the data analysis and iterate analysis parameters.

**Real-Time XFEL Data Analysis at SLAC and NERSC: a Trial Run of Nascent Exascale Experimental Data Analysis.** DOI:[10.48550/arXiv.2106.11469](https://doi.org/10.48550/arXiv.2106.11469)

Johannes P. Blaschke, Aaron S. Brewster, Daniel W. Paley, Derek Mendez, Asmit Bhowmick, Nicholas K. Sauter, Wilko Kröger, Murali Shankar, Bjoern Enders, Deborah Bard

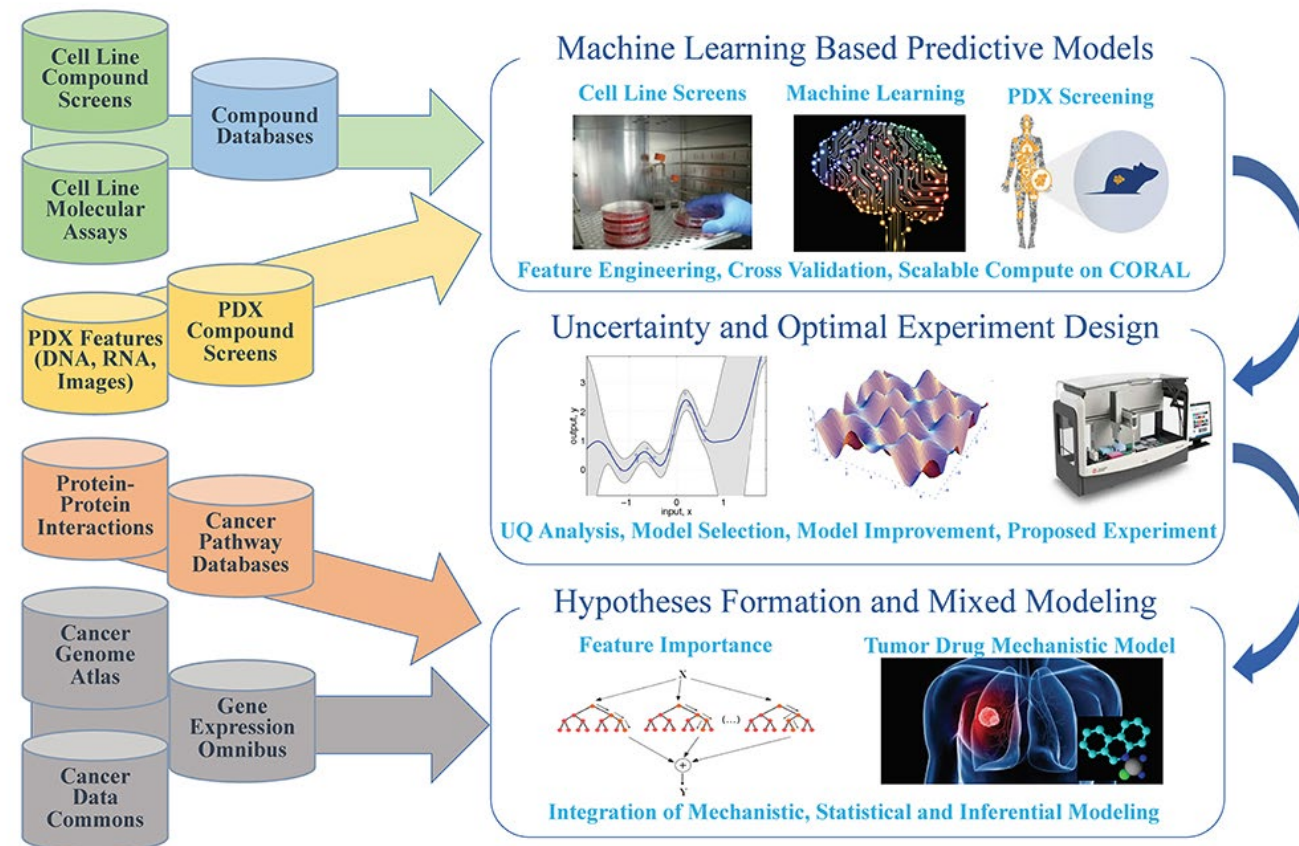


# Common scientific workflow patterns

## Hybrid

- Combines elements of multiple workflow patterns

## Predictive Models for Pre-Clinical Screening



Bhattacharya Tanmoy, Brettin Thomas, Doroshow James H., Evrard Yvonne A., Greenspan Emily J., Gryshuk Amy L., Hoang Thuc T., Lauzon Carolyn B. Veal, Nissley Dwight, Penberthy Lynne, Stahlberg Eric, Stevens Rick, Streitz Fred, Tourassi Georgia, Xia Fangfang, Zaki George, "AI Meets Exascale Computing: Advancing Cancer Research With Large-Scale High Performance Computing", Frontiers in Oncology, vol 9, 2019, DOI: [10.3389/fonc.2019.00984](https://doi.org/10.3389/fonc.2019.00984)

# Make-like (acyclic)

```
prot_dcd:
  resource:
    time: 1.0
    cpu: 4
    nrs: 1
  inp:
    - Production/remd_{n}/prot.xtc
    - Production/remd_{n}/prot.tpr
  out:
    - Production/remd_{n}/prot.psf
    - Production/remd_{n}/prot.dcd
  script: |
    . /ccs/proj/bif128/setup-env-vmd.sh
    {mpirun} vmd -dispdev text \
      -e $SRC/vmd/rewrap.tcl \
      -args {inp[0]} {inp[1]} \
        {out[0]} {out[1]}
```

## Typical characteristics:


- List of a -> b conversion rules
- File presence = workflow state
- Shell-like variable substitution
- Can be auto-parallelized

## Implementations:

GNU make  
plan9 mk  
pmake  
snakemake  
fireworks  
swift/T  
nextflow.io  
DPLASMA



# Workflow-like Programming Models: Map-Reduce




- Map: run function over all data items
- Reduce: aggregate results (sum/etc.)
- Scan: calculating running reduce
- Collect: gather info. to root rank
- Filter: conditionally remove elements
- Repartition/regroup: reorder elements

## Implementations:

module / mpi  
GNU parallel

pypi / mpi\_list  
pypi / pyspark / pysparkling  
pypi / mpi4py

# Workflow-like Execution Models: Bag of Tasks



1. Request task from server
2. Do work and output
3. Inform server of completion

Server stores a set of "tasks" and exports an API:

- Add task
- Fetch task
- Mark task complete
- May also track DAG state
- May be "push" instead of "pull"

Implementations:

pypi / parsl  
pypi / celery  
pypi / dask  
RADICAL/pilot  
Apache Airflow  
Ruby resque  
Hashicorp Nomad  
Slurm / LSF / etc.  
[github.com/frobnitzem/dwork](https://github.com/frobnitzem/dwork)

# Factors to consider

- Is it available on my system (discuss with your HPC center)
- How does it handle data provenance?
- Are there workflows that already do what I want?
- What are the computational requirements for each component/step?
  - Which need HPC?
  - Could I just use cloud vs. HPC (scale / other constraints)
- What data requirements does my problem have?
  - How much data is consumed/produced by each step?
  - Where do data transfers need to happen?

# Reproducibility challenges for HPC systems

- Library and software versions are changing constantly so it is difficult to know what versions were used at build time vs run time
  - Can be changed by the user (e.g. modules) or by system admins
  - Hotfixes for system libraries
- Kernel hardware drivers can change and affect behavior
- Files on the high-performance filesystems are "purged" using a time-based algorithm. So, as data ages, it disappears unless moved to a stable location
- Hardware may experience rare, but random system failures, terminating an otherwise correct run early or, in the worst case, silently corrupting data
- Other jobs running on the system may interfere with the network or I/O performance, leading to unpredictable amounts of latency

# Containers



- Containers are one way to reduce portability issues, however can come with a performance penalty
- Used to package up all code and dependencies necessary to execute single process or task
- Encapsulates the entire software ecosystem (minus the kernel)
- Uses a host operating system virtualization mechanism (not virtual machine)
- Common container systems: Docker (the first); Apptainer (was singularity); Podman (RedHat); Charliecloud (LANL); Shifter (NERSC); Sarus (CSCS)



# Containers (cont...)

- Originally designed for Cloud
- Docker not a good fit for HPC
  - Requires root-level access
  - Architecture does not suit most HPC installations
- Other container technologies are more tailored to HPC
- Growing HPC center support (Apptainer/Singularity, Charliecloud, Shifter, Sarus)
- May actually decrease run-time for python apps (via reduced dependency load time)
- Utilized by GitHub actions/GitLab pipelines

# Open Container Initiative (OCI)

<https://opencontainers.org>



- Industry standard for container formats and runtimes
- Allows containers built using one container system to be run with another
  - E.g. build with Podman and run using Singularity
- Containers can be stored in different registries
- Three specifications:
  - Runtime Specification (runtime-spec)
  - Image Specification (image-spec)
  - Distribution Specification (distribution-spec)

# Containers supported by HPC systems<sup>1</sup>

- ALCF
  - Theta: Singularity
  - Aurora: Singularity
- OLCF
  - Summit: Singularity
  - Frontier: Singularity
- NERSC
  - Cori: Shifter
  - Perlmutter: Shifter, Singularity, Podman
- LLNL
  - Sierra/Lassen: Singularity
  - El Capitan: Singularity
- LANL
  - Trinity: Charliecloud
  - Linux clusters: Charliecloud
  - Crossroads: Charliecloud
- SNL
  - Astra: Singularity, Charliecloud, Podman
  - Linux clusters: Shifter, Singularity

<sup>1</sup> Source: Containers and the Truth between HPC & Cloud System Software Convergence. 2021.  
DOI: [10.2172/1859696](https://doi.org/10.2172/1859696)



# Issues using containers in HPC

- MPI with HPC containers is a mess
  - Certain vendors don't support MPI libraries via a container
  - Users have to match MPI build to host MPI configurations
  - There are solutions but at the expense of portability and specific to implementation
- Build environment does not match run environment
  - Supporting non-x86 machines (eg: Summit, Astra, ...) requires special build nodes
  - Even x86 is no guarantee of performance or portability
  - Many libraries make build-time assumptions which can be bad
- Job life is a batch in HPC
  - Containers & MPI libraries don't default to playing well with schedulers
  - Module settings (environment variables) can leach into containers

# Issues using containers in HPC (cont...)

- Most HPC Container runtimes support GPUs
  - `singularity run --nv ...`
  - `shifter --module=cuda-10.1 ...`
  - `ch-fromhost --nvidia ...`
- Variation between GPUs
  - `singularity run --nv ...`
  - `singularity run --rocm ...`
- Relatively little integration to date with Intel GPUs
  - E4S image with OneAPI is planned for release next year
- None of this looks like the Cloud

# Summary

- Workflow systems provide a powerful tool to undertake scientific discovery
  - High-performance Simulation and Modeling; High performance AI; Scientific Data Lifecycle; Real-Time/Streaming; Hybrid
- There are a large number of systems available
  - Choosing the right one depends on the needs of the application (computational, data, reproducibility) and what is available locally
- Workflows help with reproducibility, there are resources that help with sharing workflow codes and data
- HPC has many reproducibility issues
  - Containers help with portability, and most HPC systems support containers
  - HPC and containers are not currently a great fit

# Other resources

## Workflow systems

- AiiDA: <https://www.aiida.net>
- BEE: <https://github.com/lanl/BEE>
- COMPSs: <https://compss.bsc.es>
- Covalent: <https://www.covalent.xyz>
- Cromwell: <http://cromwell.readthedocs.io>
- FireWorks: <https://materialsproject.github.io/fireworks>
- Galaxy: <https://galaxyproject.org>
- Maestro: <https://maestrowf.readthedocs.io>
- Nextflow: <https://www.nextflow.io>
- Pegasus: <https://pegasus.isi.edu>
- Snakemake: <https://snakemake.github.io>
- Swift: <http://swift-lang.org/Swift-T>
- Taskvine: <https://ccl.cse.nd.edu/software/taskvine>

## Containers

- Apptainer (singularity) <https://apptainer.org>
- Charliecloud: <https://hpc.github.io/charliecloud>
- Docker: <https://docker.com>
- Podman: <https://podman.io>
- Sarus: <https://sarus.readthedocs.io/en/stable>
- Shifter: <https://shifter.readthedocs.io/en/latest>

# Other resources

## Workflow repositories

- @nf-core (Nextflow)
  - <https://nf-co.re>
- Snakemake workflow catalog
  - <https://snakemake.github.io/snakemake-workflow-catalog>

## Metadata frameworks

- Common workflow language
  - <https://www.commonwl.org>
- Workflow RO-Crate
  - <https://w3id.org/workflowhub/workflow-ro-crate/1.0>
- Bioschemas Profiles
  - <https://bioschemas.org/profiles>

## Workflow registries

- WorkflowHub
  - <https://workflowhub.eu>
- Dockstore
  - <https://dockstore.org>

## Data repositories

- Zenodo
  - <https://zenodo.org>
- Dataverse
  - <https://dataverse.org>

# Other resources

- C. Harris, P. O'Leary, M. Grauer, A. Chaudhary, C. Kotfila and R. O'Bara, "Dynamic Provisioning and Execution of HPC Workflows Using Python," 2016 6th Workshop on Python for High-Performance and Scientific Computing (PyHPC), Salt Lake City, UT, USA, 2016, pp. 1-8, doi: [10.1109/PyHPC.2016.005](https://doi.org/10.1109/PyHPC.2016.005).
- S. Bharathi, A. Chervenak, E. Deelman, G. Mehta, M. -H. Su and K. Vahi, "Characterization of scientific workflows," *2008 Third Workshop on Workflows in Support of Large-Scale Science*, Austin, TX, USA, 2008, pp. 1-10, doi: [10.1109/WORKS.2008.4723958](https://doi.org/10.1109/WORKS.2008.4723958).
- Blaschke, Johannes P., Brewster, Aaron S., et. al. 2021. "Real-Time XFEL Data Analysis at SLAC and NERSC: a Trial Run of Nascent Exascale Experimental Data Analysis". doi: [10.48550/arXiv.2106.11469](https://doi.org/10.48550/arXiv.2106.11469).
- Younge, Andrew J. 2021. "Containers and the Truth between HPC & Cloud System Software Convergence.". DOI: [10.2172/1859696](https://doi.org/10.2172/1859696).