



# Spack: Package Management for HPC

ATPESC 2024 Software Track

Todd Gamblin  
Spack Project Lead  
Distinguished Member of Technical Staff  
Lawrence Livermore National Laboratory

THE **LINUX** FOUNDATION



# Background

THE **LINUX** FOUNDATION



# We build codes from hundreds of small, complex pieces

*Just when we're starting to solve the problem of how to create software using reusable parts, it founders on the nuts-and-bolts problems outside the software itself.*

P. DuBois & T. Epperly. ***Why Johnny Can't Build***. Scientific Programming. Sep/Oct 2003.

- Component-based software development dates back to the 60's
  - M.D. McIlroy, *Mass Produced Software Components*. NATO SE Conf., 1968
- **Pros are well known:**
  - Teams can and must reuse each others' work
  - Teams write less code, meet deliverables faster
- **Cons:**
  - Teams must ensure that components work together
  - Integration burden increases with each additional library
  - Integration must be repeated with each update to components
  - **Components must be vetted!**
- **Managing changes over time is becoming intractable**



Build-time incompatibility; fail fast



Appears to work; subtle errors later

# Modern scientific codes rely on icebergs of dependency libraries

**71 packages**  
**88 dependencies**

## LBANN: Neural Nets for HPC

## MFEM: Higher-order finite elements

**31 packages,  
69 dependencies**

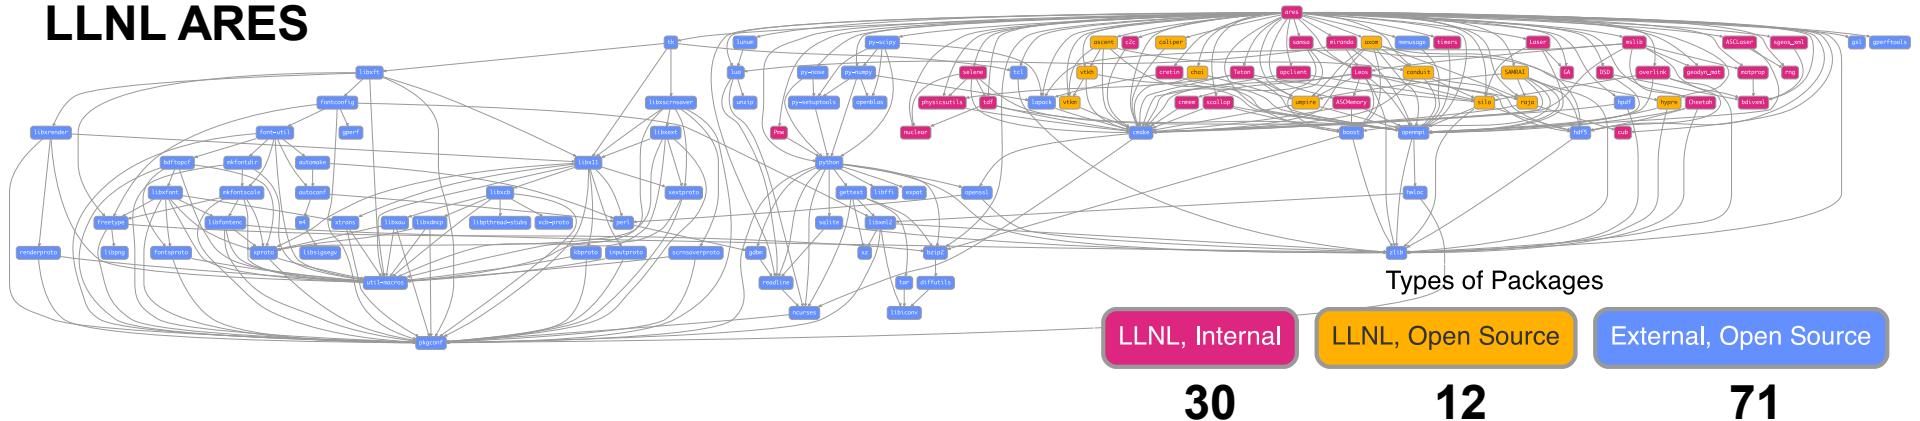
# R Genome Data Analysis Tools

**179 packages,**  
**527 dependencies**

**179 packages,  
527 dependencies**

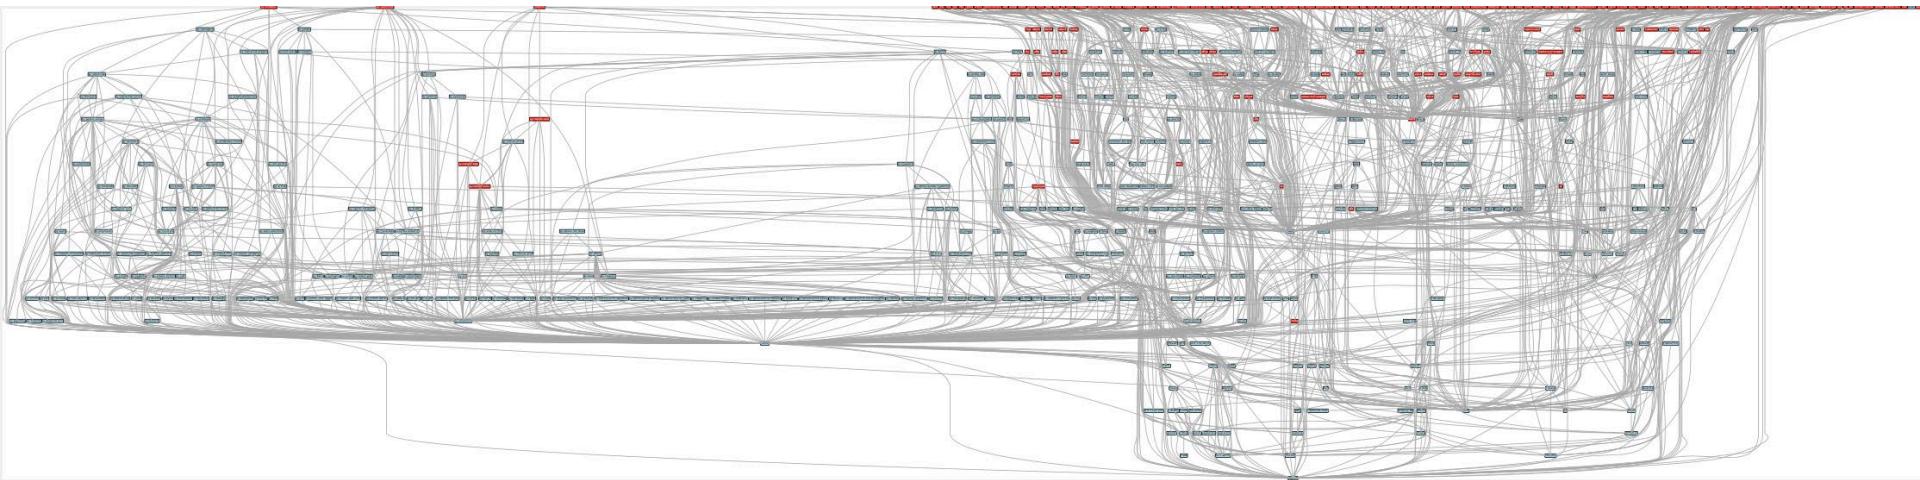
# Even proprietary software builds on top of open source

## LLNL ARES



- Open source is critical for nearly every application
- We *cannot* replace all these OSS components with our own
  - How do we put them all together effectively?
  - Do you *have* to integrate this stuff by hand?

# ECP's E4S stack is even larger than these codes

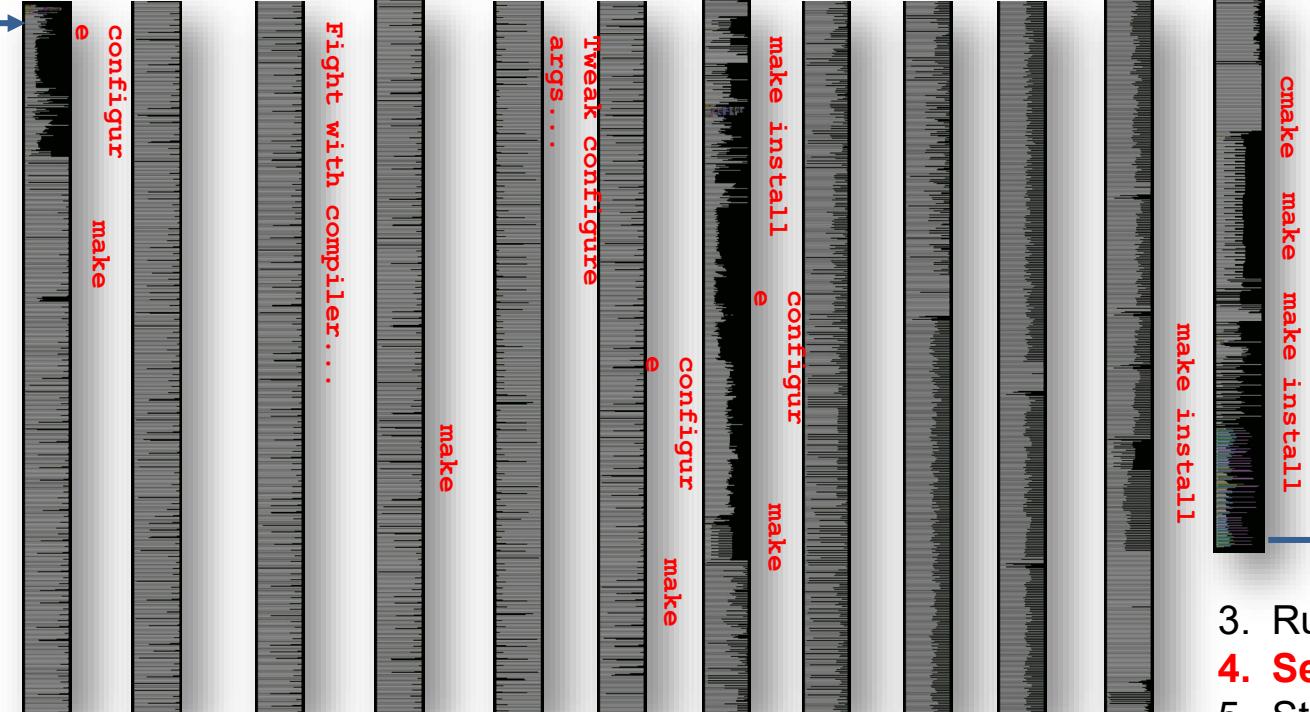


- Red boxes are the packages in it (about 100)
- Blue boxes are what *else* you need to build it (about 600)
- It's infeasible to build and integrate all of this manually

# Some history:

## How to install software on a supercomputer, circa 2013

1. Download all 16 tarballs you need
2. Start building!



3. Run code
4. Segfault!?
5. Start over...

spack.io



# What does a package managers do to help?

- Does not replace Cmake/Autotools
- Manages ***dependencies***
  - Drives package-level build systems
  - Ensures **consistency** and **compatibility** among builds of packages in the ecosystem
- Stores **community knowledge**
  - Cache of package build recipes
  - Determining magic configure lines takes time

Package Manager

- Package installation
- Dependency relationships, conflicts
- May drive package-level build systems

High Level Build System

- Cmake, Autotools
- Handle library abstractions
- Generate Makefiles, etc.

Low Level Build System

- Make, Ninja
- Handles dependencies among *commands* in a single build

# Many package managers (conda, pip, apt, etc.) make simplifying assumptions about the software ecosystem

---

- **1:1 relationship between source code and binary (per platform)**
  - Good for reproducibility (e.g., Debian)
  - Bad for performance optimization
- **Binaries should be as portable as possible**
  - What most distributions do
  - Again, bad for performance
- **Toolchain is the same across the ecosystem**
  - One compiler, one set of runtime libraries
  - Or, no compiler (for interpreted languages) and *just one language*

Outside these boundaries, users are typically on their own

# High Performance Computing (HPC) violates many of these assumptions

- **Code is typically distributed as source**
  - With exception of vendor libraries, compilers
- **Often build many variants of the same package**
  - Developers' builds may be very different
  - Many first-time builds when machines are new
- **Code is optimized for the processor and GPU**
  - Must make effective use of the hardware
  - Can make 10-100x perf difference
- **Rely heavily on system packages**
  - Need to use optimized libraries that come with machines
  - Need to use host GPU libraries and network
- **Multi-language**
  - C, C++, Fortran, Python, others all in the same ecosystem

## Some Supercomputers



Summit  
Oak Ridge National Lab  
Power9 / NVIDIA



Fugaku  
RIKEN  
Fujitsu/ARM a64fx



Perlmutter  
Lawrence Berkeley National Lab  
AMD Zen / NVIDIA



Aurora  
Argonne National Lab  
Intel Xeon / Xe



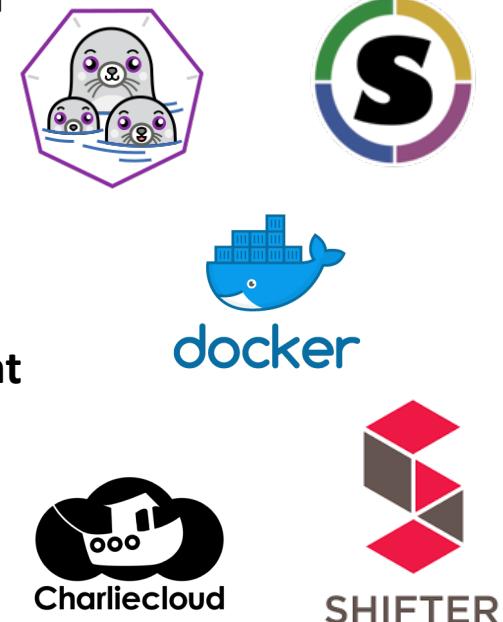
FRONTIER  
Oak Ridge National Lab  
AMD Zen / Radeon



El Capitan  
Lawrence Livermore National Lab  
AMD Zen / Radeon

# What about containers?

- Containers provide a great way to reproduce and distribute an already-built software stack
- Someone needs to build the container!
  - This isn't trivial
  - Containerized applications still have hundreds of dependencies
- Using the OS package manager inside a container is insufficient
  - Most binaries are built unoptimized
  - Generic binaries, not optimized for specific architectures
- HPC containers are often optimized per-system
  - Not clear that we can ever build one container for all facilities



We need something more flexible to **build** versions of containers

# Overview & Community

THE **LINUX** FOUNDATION



# Spack enables Software distribution for HPC

- Spack automates the build and installation of scientific software
- Packages are *parameterized*, so that users can easily tweak and tune configuration

## No installation required: clone and go

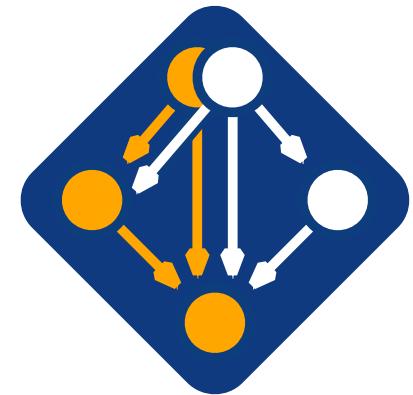
```
$ git clone https://github.com/spack/spack  
$ spack install hdf5
```

## Simple syntax enables complex installs

```
$ spack install hdf5@1.10.5  
$ spack install hdf5@1.10.5 %clang@6.0  
$ spack install hdf5@1.10.5 +threadsafe
```

```
$ spack install hdf5@1.10.5 cppflags="-O3 -g3"  
$ spack install hdf5@1.10.5 target=haswell  
$ spack install hdf5@1.10.5 +mpi ^mpich@3.2
```

- Ease of use of mainstream tools, with flexibility needed for HPC
- In addition to CLI, Spack also:
  - Generates (but does **not** require) *modules*
  - Allows conda/virtualenv-like *environments*
  - Provides many devops features (CI, container generation, more)
  - Supports *binary “buildcaches”* so that you don’t have to build everything from source



[github.com/spack/spack](https://github.com/spack/spack)

# Who can use Spack?

---

**People who want to use or distribute software for HPC!**

**1. End Users of HPC Software**

- Install and run HPC applications and tools

**2. HPC Application Teams**

- Manage third-party dependency libraries

**3. Package Developers**

- People who want to package their own software for distribution

**4. User support teams at HPC Centers**

- People who deploy software for users at large HPC sites



# Spack was critical for ECP's mission to create a robust, capable exascale software ecosystem.



The Extreme-scale Scientific Software Stack (E4S) is a community effort to provide open source software packages for developing, deploying and running scientific applications on high-performance computing (HPC) platforms. E4S provides from-source builds and containers of a broad collection of HPC software packages.

**Purpose**  
E4S aims to accelerate the development, deployment and use of HPC software, lowering the barriers for HPC users. E4S provides containers and turn-key, from-source build of more than 80 popular HPC products in programming models such as Fortran, C/C++, development tools such as PETSc, Trilinos, and MPI, math libraries such as PETSc and Trilinos, and Data and Viz tools such as HDF5 and Paraview.

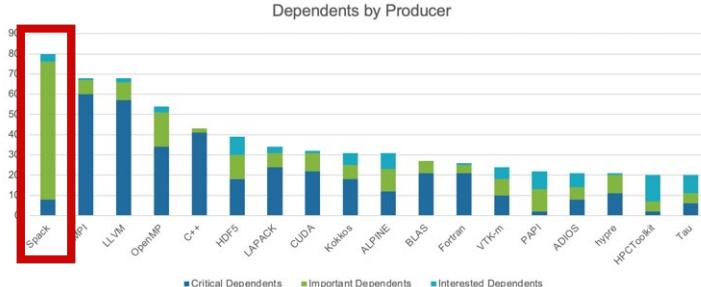
**Approach**  
By using Spack as the meta-build tool and growing containers of pre-built binaries for Docker, Singularity, Shifter and ContainerCloud, E4S enables the flexible use and testing of a large collection of reusable HPC software packages.

**Platforms**  
The E4S stack is designed to run on a variety of HPC platforms, including Cray, Intel, and AMD architectures.

**Testing**  
The E4S stack is rigorously tested across a wide range of HPC environments to ensure compatibility and performance.

<https://e4s.io>

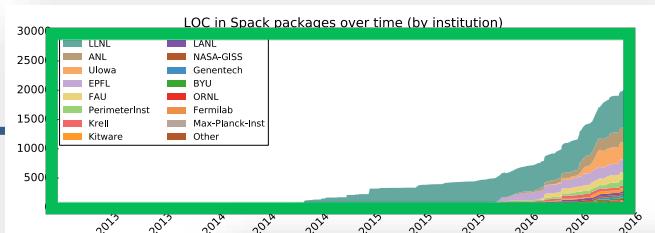
- Used for building software on the three U.S. exascale systems
- ECP built the Extreme Scale Scientific Software Stack (E4S) with Spack – more at <https://e4s.io>
- Project continues on ASC and ASCR funding



Spack was the most depended-upon project in ECP

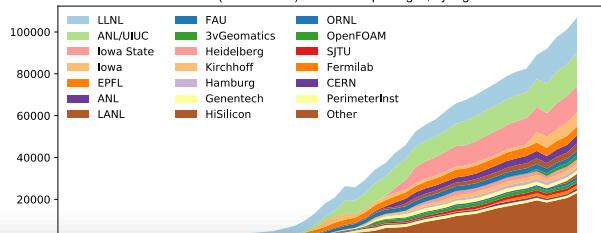
[spack.io](https://spack.io)





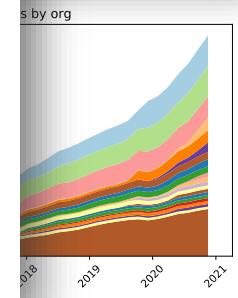
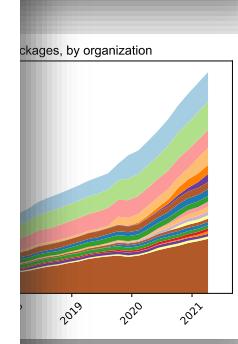
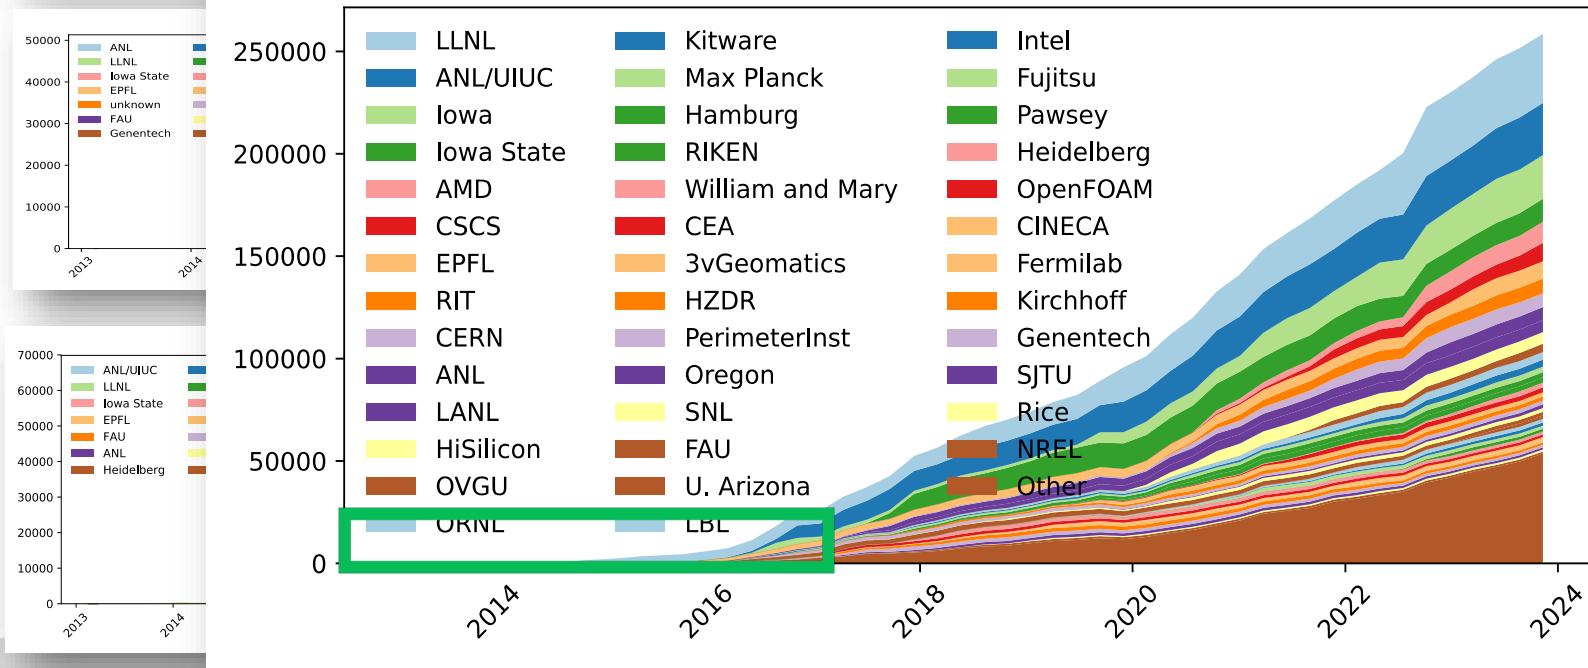
2016

Fall  
2020

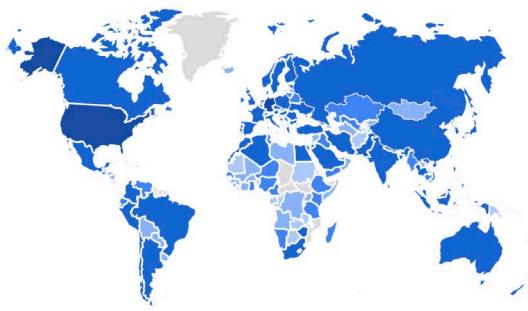


2024

Contributions (lines of code) over time in packages, by organization



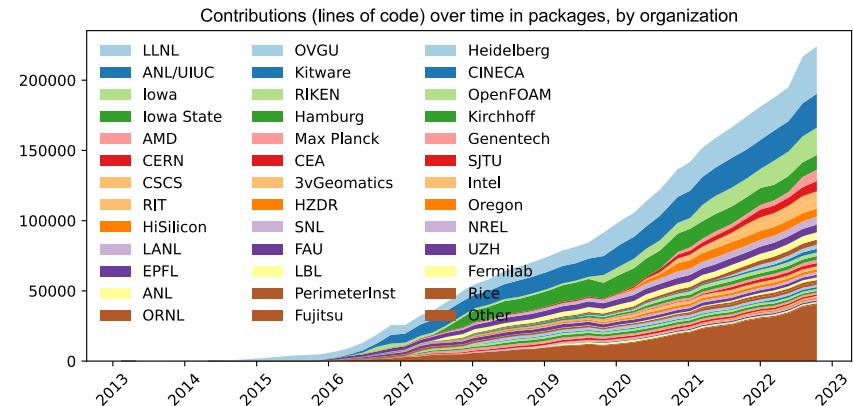
# Spack sustains the HPC software ecosystem with the help of many contributors



COUNTRY	USERS
United States	23K
Germany	5.3K
China	4.6K
India	4.5K
United Kingdom	3.3K
France	3K
Japan	2.4K

2023 aggregate documentation user counts from GA4  
(note: yearly user counts are almost certainly too large)

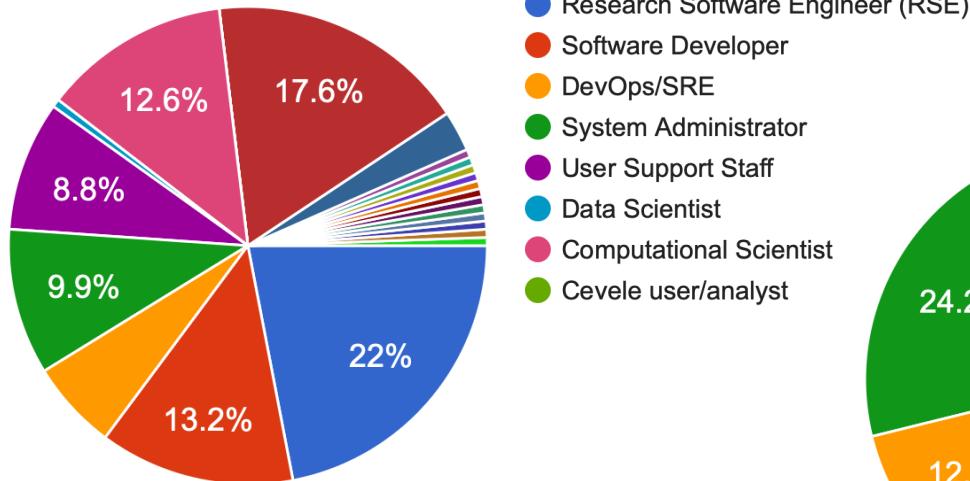
**Over 8,000 software packages**  
**Over 1,300 contributors**



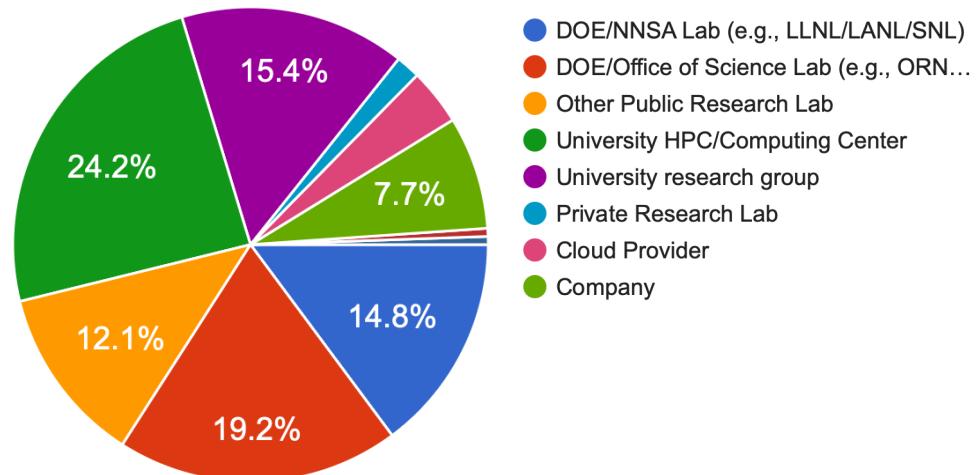
**Contributors continue to grow worldwide!**

# Spack users have diverse roles across many types of institutions

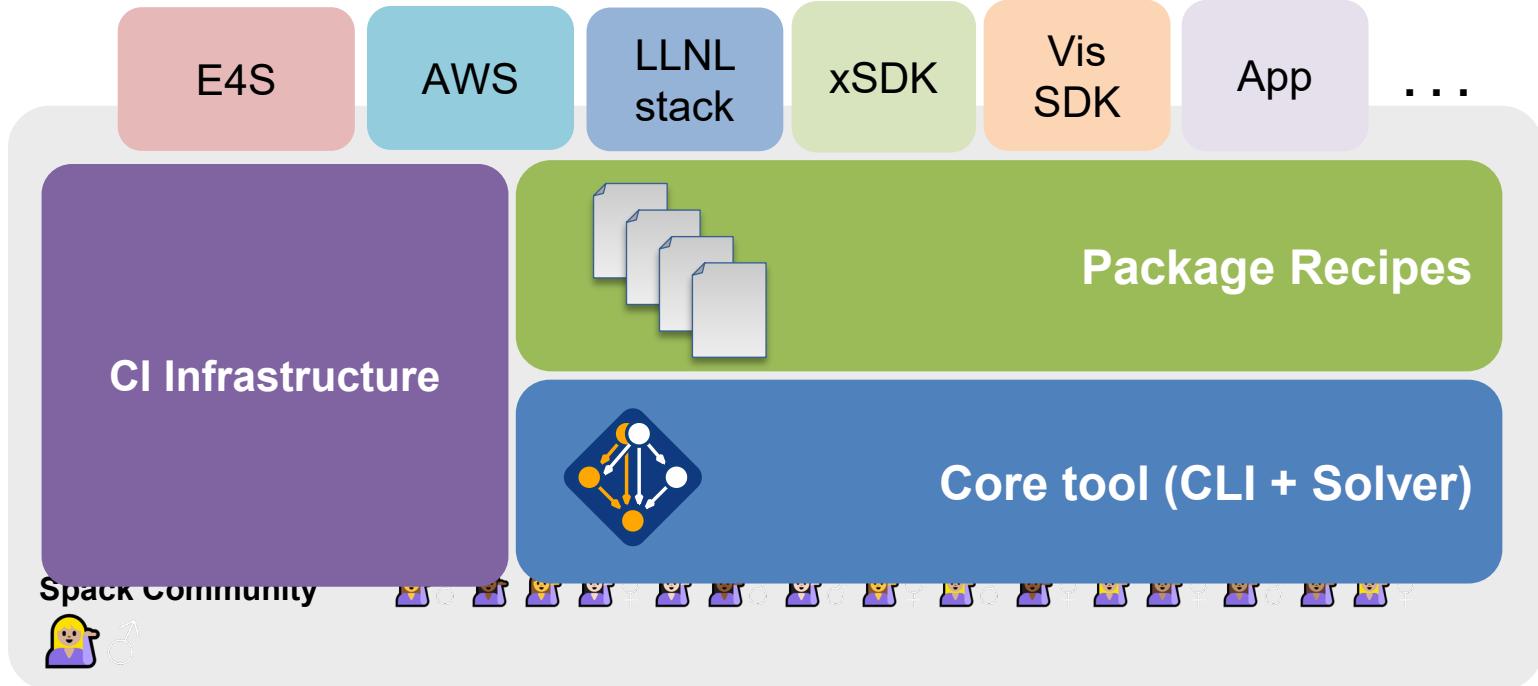
## What type of user are you?



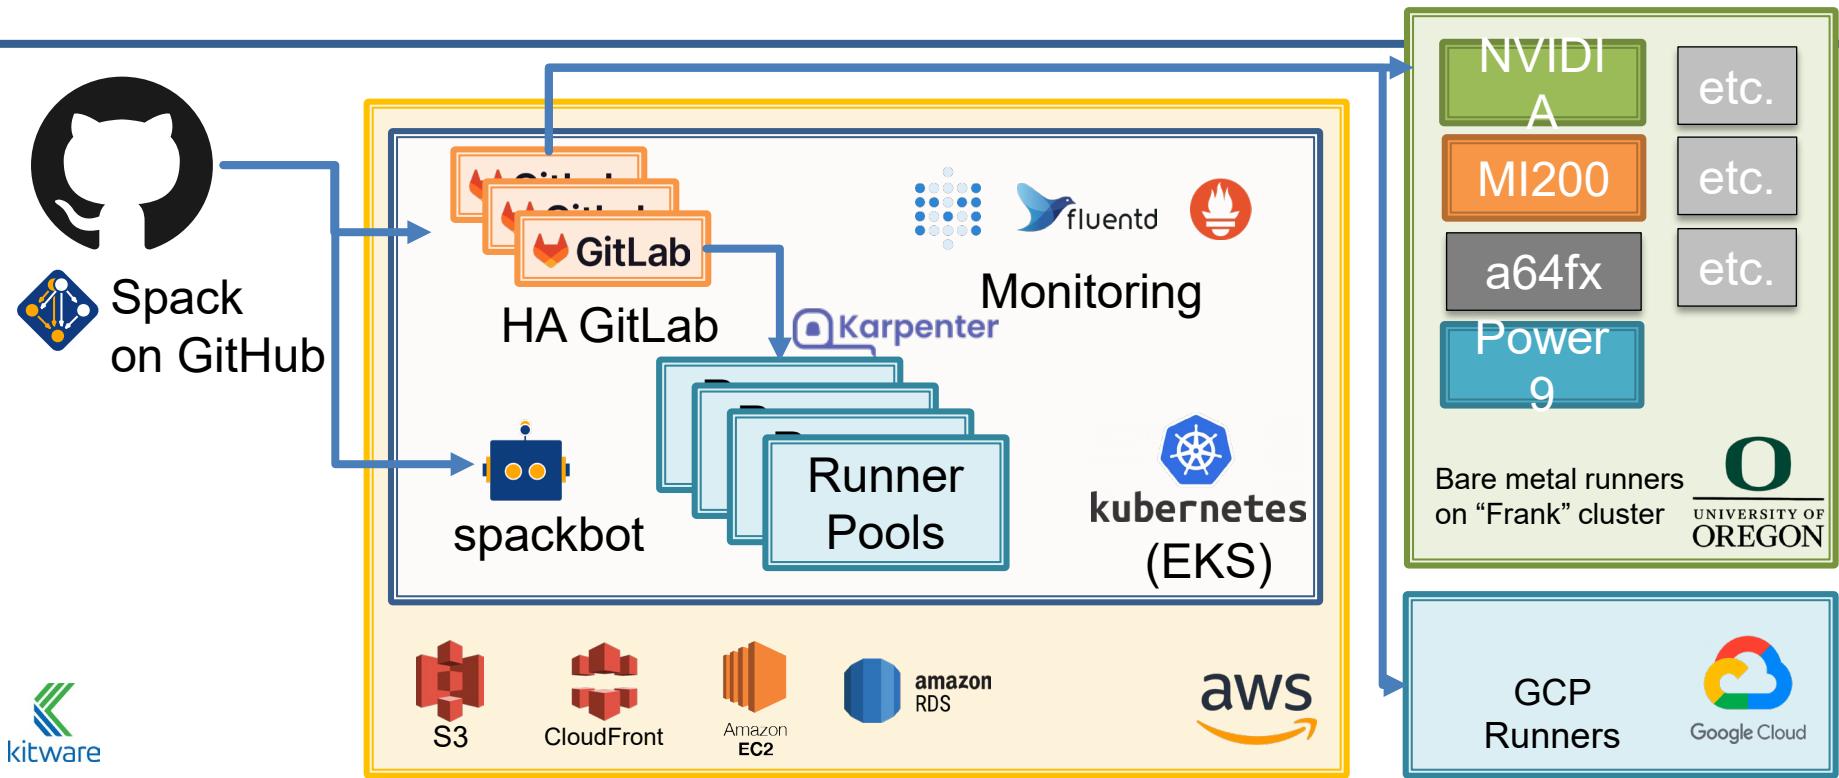
## Where do you work?



# What does the Spack project look like now?



# Spack CI Architecture



# CI enables us to maintain public *build caches*

- We test a subset of package builds:
  - In CI for every pull request
  - On release branches
  - On the develop branch
- We provide build caches for:
  - Every release
  - Latest develop branch
  - Weekly develop snapshots
  - Different architectures and software stacks

The screenshot shows the Spack Packages search interface. A search bar at the top contains the text "python". Below it, a list of packages is displayed, starting with "py-antlr4-python3-runtime" and including "py-avro", "py-biopython", "py-bx-python", "py-dnspython", "py-domdf-python-tools", and "py-edfplus-python". To the right of the search results, a "Get Started" section provides instructions for adding a build cache to a Spack configuration:

```
$ spack mirror add v0.22.1-data-vis-sdk https://binaries.spack.io/v0.22.1/data-vis-sdk
$ spack buildcache keys --install --trust
```

Below this, a table lists available builds for the "adios2" package, showing columns for UID, TAG, PACKAGE, VERSIONS, PLATFORM, OS, TARGET, and COMPILER. The table includes entries for "aarch64", "neoverse\_v1", "neoverse\_v2", "ppc64le", "x86\_64\_v3", and "x86\_64\_v4". Another table for "alsa-lib" and a third for "apcomp" and "argobots" are also partially visible.

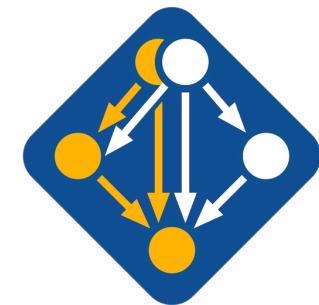
See available builds  
at [cache.spack.io](https://cache.spack.io)

Find available packages  
at [packages.spack.io](https://packages.spack.io)

# Spack recently became a Linux Foundation Project!

## What does that mean?

- Spack has a legal entity: a 501(c)(6) non-profit company
  - This is a neutral legal entity
  - Can be in legal agreements (e.g. for distributing binaries)
  - Can get discounts on, e.g., Slack
- Spack will have a Technical Steering Committee (TSC)
  - Plan is to make the main developer meetings more public
  - Also have official steering committee meetings
  - Working on initial GOVERNANCE.md, initial TSC members
- *Trademark* (Spack name, logo) assigned to Linux Foundation
- Spack project resources owned by Linux Foundation:
  - [spack.io](http://spack.io) website
  - GitHub Organization



Why are we doing this?

# Spack is not the only HPC/AI/data science package manager out there



## 1. “Functional” Package Managers

- Nix
- Guix

<https://nixos.org>  
<https://hpc.guix.info>



## 2. Build-from-source Package Managers

- Homebrew, LinuxBrew
- MacPorts
- Gentoo

<https://brew.sh>  
<https://www.macports.org>  
<https://gentoo.org>

## Other tools in the HPC Space:



### ▪ Easybuild

- An installation tool for HPC
- Focused on HPC system administrators – different package model from Spack
- Relies on a fixed software stack – harder to tweak recipes for experimentation

<https://easybuild.io>



### ▪ Conda / Mamba / Pixi

- Very popular binary package ecosystem for data science
- Not targeted at HPC; generally has unoptimized binaries

<https://conda.io>  
<https://mamba.readthedocs.io>  
<https://prefix.dev>

# Spack Usage

THE **LINUX** FOUNDATION



# Spack provides a *spec* syntax to describe customized installations

```
$ spack install mpileaks           unconstrained
$ spack install mpileaks@3.3       @ custom version
$ spack install mpileaks@3.3 %gcc@4.7.3    % custom compiler
$ spack install mpileaks@3.3 %gcc@4.7.3 +threads  +/- build option
$ spack install mpileaks@3.3 cppflags="-O3 -g3"   set compiler flags
$ spack install mpileaks@3.3 target=zen2        set target microarchitecture
$ spack install mpileaks@3.3 ^mpich@3.2 %gcc@4.9.3  ^ dependency information
```

- Each expression is a *spec* for a particular configuration
  - Each clause adds a constraint to the spec
  - Constraints are optional – specify only what you need.
  - Customize install on the command line!
- Spec syntax is recursive
  - Full control over the combinatorial build space

# Spack packages are *templates*

They use a simple Python DSL to define how to build

```
from spack import *

class Kripke(CMakePackage):
    """Kripke is a simple, scalable, 3D Sn deterministic particle
       transport proxy/mini app.

    homepage = "https://computation.llnl.gov/projects/co-design/kripke"
    url     = "https://computation.llnl.gov/projects/co-design/download/kripke-openmp-1.1.tar.gz"

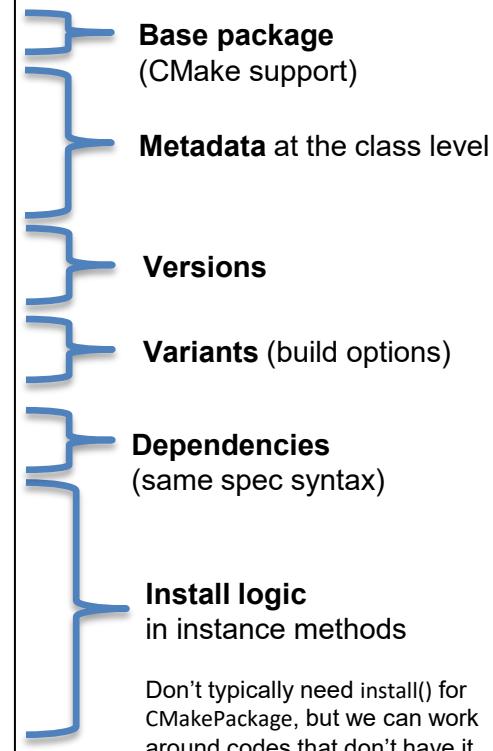
    version('1.2.3', sha256='3f7f2eef0d1ba5825780d626741eb0b3f026a096048d7ec4794d2a7dfbe2b8a6')
    version('1.2.2', sha256='eaf9ddf562416974157b34d00c3a1c880fc5296fce2aa2efa039a86e0976f3a3')
    version('1.1', sha256='232d74072fc7b848fa2adc8a1bc839ae8fb5f96d50224186601f55554a25f64a')

    variant('mpi', default=True, description='Build with MPI.')
    variant('openmp', default=True, description='Build with OpenMP enabled.')

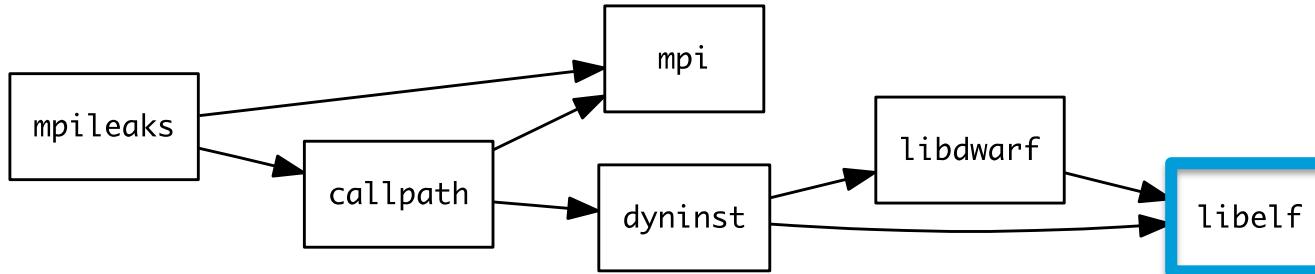
    depends_on('mpi', when='+mpi')
    depends_on('cmake@3.0:', type='build')

    def cmake_args(self):
        return [
            '-DENABLE_OPENMP=%s' % ('+openmp' in self.spec),
            '-DENABLE_MPI=%s' % ('+mpi' in self.spec),
        ]

    def install(self, spec, prefix):
        # Kripke does not provide install target, so we have to copy
        # things into place.
        mkdirp(prefix.bin)
        install('../spack-build/kripke', prefix.bin)
```



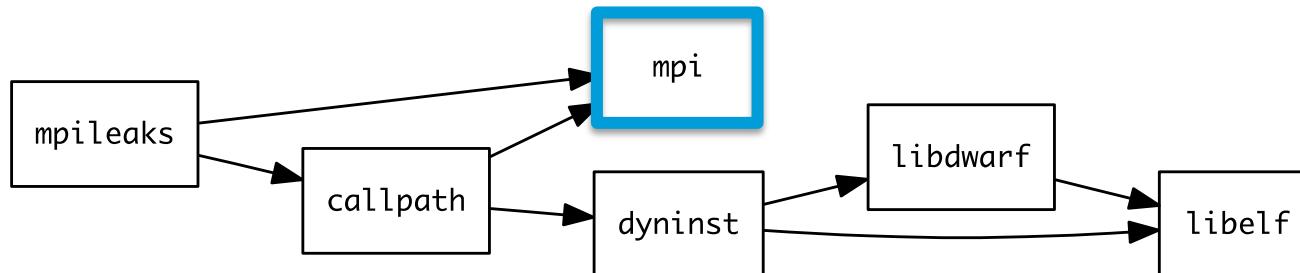
# Spack Specs can constrain versions of dependencies



```
$ spack install mpileaks %intel@12.1 ^libelf@0.8.12
```

- Spack ensures *one* configuration of each library per DAG
  - Ensures ABI consistency.
  - User does not need to know DAG structure; only the dependency *names*.
- Spack can ensure that builds use the same compiler, or you can mix
  - Working on ensuring ABI compatibility when compilers are mixed.

# Spack handles ABI-incompatible, versioned interfaces like MPI



- *mpi* is a *virtual dependency*
- Install the same package built with two different MPI implementations:

```
$ spack install mpileaks ^mvapich@1.9
```

```
$ spack install mpileaks ^openmpi@1.4:
```

- Let Spack choose MPI implementation, as long as it provides MPI 2 interface:

```
$ spack install mpileaks ^mpi@2
```

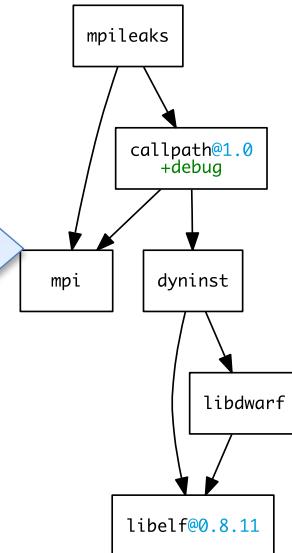
# Concretization fills in missing configuration details when the user is not explicit.

mpileaks ^callpath@1.0+debug ^libelf@0.8.11

User input: abstract spec with some constraints

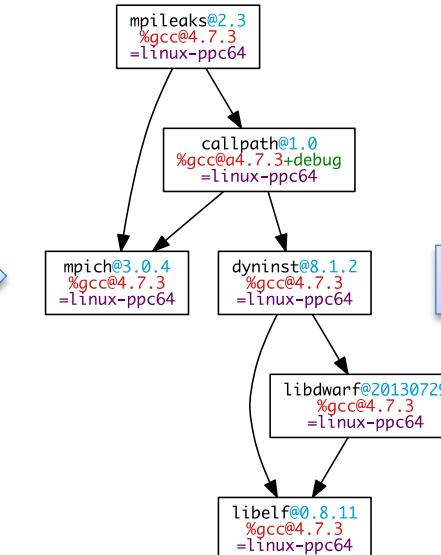
spec.yaml

Normalize



Abstract, normalized spec with some dependencies

Concretize



Concrete spec is fully constrained and can be passed to install

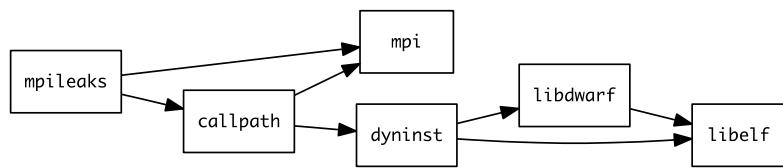
Store

```
spec:  
- mpileaks:  
  arch: linux-x86_64  
  compiler:  
    name: gcc  
    version: 4.9.2  
  dependencies:  
    adept-utils: ksrtkpbzac3ss2ixcjkcorlabybnptp4  
    callpath: bah5f4h4d2n47mgycej2mtrnrivxy72t3  
    mpich: aa4ar6ifj23yijqmdabeakpejcli72t3  
    hash: 33hjhix7p6gvyzn5ptgyes7ghyprujh  
    variants: {}  
    version: '1.0'  
- adept-utils:  
  arch: linux-x86_64  
  compiler:  
    name: gcc  
    version: 4.9.2  
  dependencies:  
    boost: teesjv7ehpe5ksspjm5dk43a7qnowlq  
    mpich: aa4ar6ifj23yijqmdabeakpejcli72t3  
    hash: ksrtkpbzac3ss2ixcjkcorlabybnptp4  
    variants: {}  
    version: 1.0.1  
- boost:  
  arch: linux-x86_64  
  compiler:  
    name: gcc  
    version: 4.9.2  
  dependencies: {}  
  hash: teesjv7ehpe5ksspjm5dk43a7qnowlq  
  variants: {}  
  version: 1.59.0  
...
```

Detailed provenance stored with installed package

# Spack handles combinatorial software complexity

## Dependency DAG



## Installation Layout



```
opt
└── spack
    ├── linux-rhel7-skylake
    │   └── gcc-8.3.0
    │       ├── mpileaks-1.0-hc4sm4vuzpm4znmvrfzri4ow2mkphe2e
    │       ├── callpath-1.0.4-daqqpssxb6qbfrztsezkmhus3xoflpsy
    │       ├── openmpi-4.1.4-u64v26igvxxy23hysmklfums6tgjv5r
    │       ├── dyninst-12.1.0-u64v26igvxyn23hysmklfums6tgjv5r
    │       ├── libdwarf-20180129-u5eawkvaoc7vonabe6nndkcfwuv233cj
    │       └── libelf-0.8.13-x46q4wm46ay4pltrijbgizxjrbaka6
```

- Each unique dependency graph is a unique **configuration**.
- Each configuration in a unique directory.
  - Multiple configurations of the same package can coexist.
- **Hash** of entire directed acyclic graph (DAG) is appended to each prefix.
- Installed packages automatically find dependencies
  - Spack embeds RPATHs in binaries.
  - No need to use modules or set LD\_LIBRARY\_PATH
  - Things work *the way you built them*

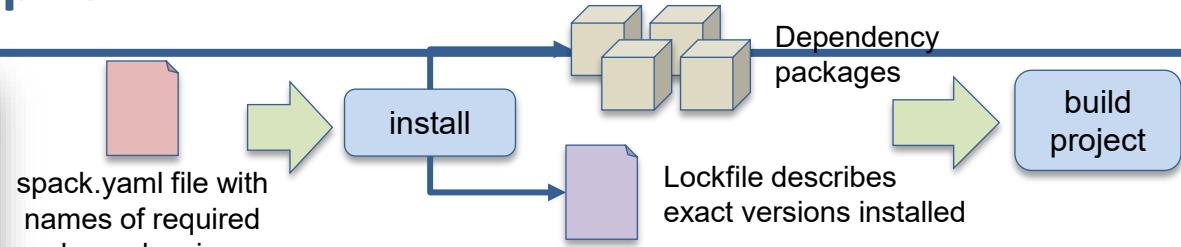
# Spack environments enable users to build customized stacks from an abstract description

## Simple spack.yaml file

```
spack:  
  # include external configuration  
  include:  
    - ./special-config-directory/  
    - ./config-file.yaml  
  
  # add package specs to the `specs` list  
  specs:  
    - hdf5  
    - libelf  
    - openmpi
```

## Concrete spack.lock file (generated)

```
{  
  "concrete_specs": {  
    "6s63so2kstp3zyvjezglndmavy6l3nul": {  
      "hdf5": {  
        "version": "1.10.5",  
        "arch": {  
          "platform": "darwin",  
          "platform_os": "mojave",  
          "target": "x86_64"  
        },  
        "compiler": {  
          "name": "clang",  
          "version": "10.0.0-apple"  
        },  
        "namespace": "builtin",  
        "parameters": {}  
      }  
    }  
  }  
}
```



- spack.yaml describes project requirements
- spack.lock describes exactly what versions/configurations were installed, allows them to be reproduced.
- Can also be used to maintain configuration together with Spack packages.
  - E.g., versioning your own local software stack with consistent compilers/MPI implementations
  - Allows developers and site support engineers to easily version Spack configurations in a repository

**Environments have enabled us to add build many features to support developer workflows**

```

class Cmake(Package):
    executables = ['cmake']

    @classmethod
    def determine_spec_details(cls, prefix, exes_in_prefix):
        exe_to_path = dict(
            (os.path.basename(p), p) for p in exes_in_prefix
        )
        if 'cmake' not in exe_to_path:
            return None

        cmake = spack.util.executable.Executable(exe_to_path['cmake'])
        output = cmake('--version', output=str)
        if output:
            match = re.search(r'cmake\.\w+version\s+(\S+)', output)
            if match:
                version_str = match.group(1)
                return Spec('cmake@{0}'.format(version_str))

```

## package.py

## spack.yaml configuration

```
spaceman
Automate
packages:
cmake:
externals:
- spec: cmake@3.15.1
  prefix: /usr/local
```

## spack external find

Automatically find and configure external packages on the system

## spack.yaml

## spack ci

Automatically generate parallel build pipelines  
(more on this later)

gitlab-ci.yml CI pipeline

## **spack containerize**

## Turn environments into container build recipes

```

class LibSigsegv(AutoToolPackage):
    """GNU libsigsegv is a library for handling page faults in user mode.

    # ... package contents ...

    extra_install = ['tests/libsigsegv']

    def test(self):
        data_dir = self.test_suite.current_test_data_dir
        smoke_test_c = data_dir.join('smoke_test.c')

        self.run_test(
            'cc', [
                '-I', data_dir.parent.path('include'),
                '-L', data_dir.parent.path('lib'),
                '-lsigsegv',
                smoke_test_c,
                '-o', 'smoke_test'
            ],
            purpose='check linking'
        )

        self.run_test(
            'smoke_test', [], data_dir.join('smoke_test.out'),
            purpose='run smoke test'
        )

    def run_test(self, *args, **kwargs):
        """Run a command and capture its output.

        :param list args: command arguments
        :param dict kwargs: keyword arguments
        :return: tuple (exit_code, output)
        """
        if 'purpose' in kwargs:
            purpose = kwargs['purpose']
            if purpose == 'check linking':
                self._check_linking(*args)
            elif purpose == 'run smoke test':
                self._run_smoke_test(*args)
            else:
                raise ValueError("Unknown purpose: %s" % purpose)
        else:
            return super().run_test(*args, **kwargs)

    def _check_linking(self, *args):
        """Check if the command links correctly.

        :param list args: command arguments
        """
        exit_code, output = self._run(*args)
        if exit_code != 0:
            raise RuntimeError("Linking failed: %s" % output)

    def _run_smoke_test(self, *args):
        """Run the smoke test and check its output.

        :param list args: command arguments
        """
        exit_code, output = self._run(*args)
        if exit_code != 0:
            raise RuntimeError("Smoke test failed: %s" % output)

```

## package.py



# spack develop lets developers work on many packages at once

- Developer features so far have focused on single packages (spack dev-build, etc.)
- New spack develop feature enables development environments
  - Work on a code
  - Develop multiple packages from its dependencies
  - Easily rebuild with changes
- Builds on spack environments
  - Required changes to the installation model for dev packages
  - dev packages don't change paths with configuration changes
  - Allows devs to iterate on builds quickly

```
$ spack env activate .
$ spack add myapplication
$ spack develop axom@0.4.0
$ spack develop mfem@4.2.0

$ ls
spack.yaml      axom/      mfem/

$ cat spack.yaml
spack:
  specs:
    - myapplication      # depends on axom, mfem

  develop:
    - axom @0.4.0
    - mfem @develop
```

# Teams are building development front-ends top of on Spack

- Many different approaches:
  - Thin `mack` wrapper for MARBL team at LLNL
  - `spack-manager` at Sandia
  - `SpackDev` at Fermilab
  - `spack-organizer` at CEA
  - `spack_cmake` at LANL
- Workflow seems to be converging around environments + spack develop
- We are trying to bring many of the features from these scripts into core
  - Most of the front-ends are opinionated in one way or another
  - We want spack to support but not force these workflows

EPJ Web of Conferences **245**, 05035 (2020)  
*CHEP 2019*

FERMILAB-CONF-20-635-SCD  
<https://doi.org/10.1051/epjconf/202024505035>

## SpackDev: Multi-Package Development with Spack

Chris Green<sup>1,\*</sup>, James Amundson<sup>1</sup>, Lynn Garren<sup>1</sup>, Patrick Gartung<sup>1</sup>, and Elizabeth Sexton-Kennedy<sup>2</sup>

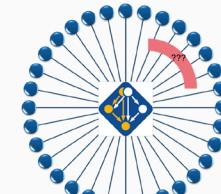
<sup>1</sup>Scientific Computing Division, Fermi National Accelerator Laboratory, Batavia, IL 60510, USA  
<sup>2</sup>Office of the Chief Information Officer, Fermi National Accelerator Laboratory, Batavia, IL 60510, USA

**Spack-Manager**

**Abstract.** High Energy Physics (HEP) projects often have hundreds of external software packages, each with its own configuration and dependency requirements. Managing coherent configurations across multiple machines can be challenging enough, but managing configurations for a large body of code can be even more difficult, especially when there is room for error. Spack is a popular package manager for HEP, but it does not focus on the needs of the end user. Spack is designed for system administrators whose primary concern is managing testing, and installations of software packages. Spack is a similarly stable dependency manager as CMake, but within HEP as that code is often developed by multiple teams. Efforts to develop future versions of Spack will focus on those platforms [1].

**Spack-Manager**

**Spack-Manager** is a light-weight extension to **Spack** that is intended to streamline the software development and deployment cycle for software projects on specific machines. A given software project typically has multiple configurations across many machines. Spack-Manager is quite literal in its name, in that it provides a way to manage and organize these configurations across multiple machines, and multiple projects.



1 Spack is serving the package management needs of thousands of software packages. However, individual application teams will likely have their own needs for Spack for their individual applications specific to their needs.

New command: `spack cmake` - Configure a CMake project using a Spack spec #45494

Conversation 0 Commits 2 Checks 31 Files changed 4

rbberger commented 3 days ago - edited

@tgambilin Here is my attempt to upstream the tool I mentioned in Slack. I've modified it a bit to be more like a regular Spack command, reworked some of the output, added a `--dry-run` option, and stripped some project specific features. The code can likely still be generalized a bit more. Right now, it is very specific to C/C++-based projects that use Kokkos for HIP and CUDA. Below is a bit of motivating text, part of which might end up in the final documentation.

Let me know what you think and how it could be improved to better fit into Spack.

This command was built out of a need to simplify developer's workflows while working on codes that also provide Spack packages. The original `spack_cmake` tool was written for the development and CI testing of FileCSE.

We recognized that many complicated use cases and the corresponding CMake command-lines were already encoded in the Spack package for a given code, and we needed a way to extract that useful logic in a pure CMake-based

Reviewers tgambilin prakewich

At least 1 approving review is required to merge this pull request.

Assignees No one assigned yourself

Labels commands core new shell-support

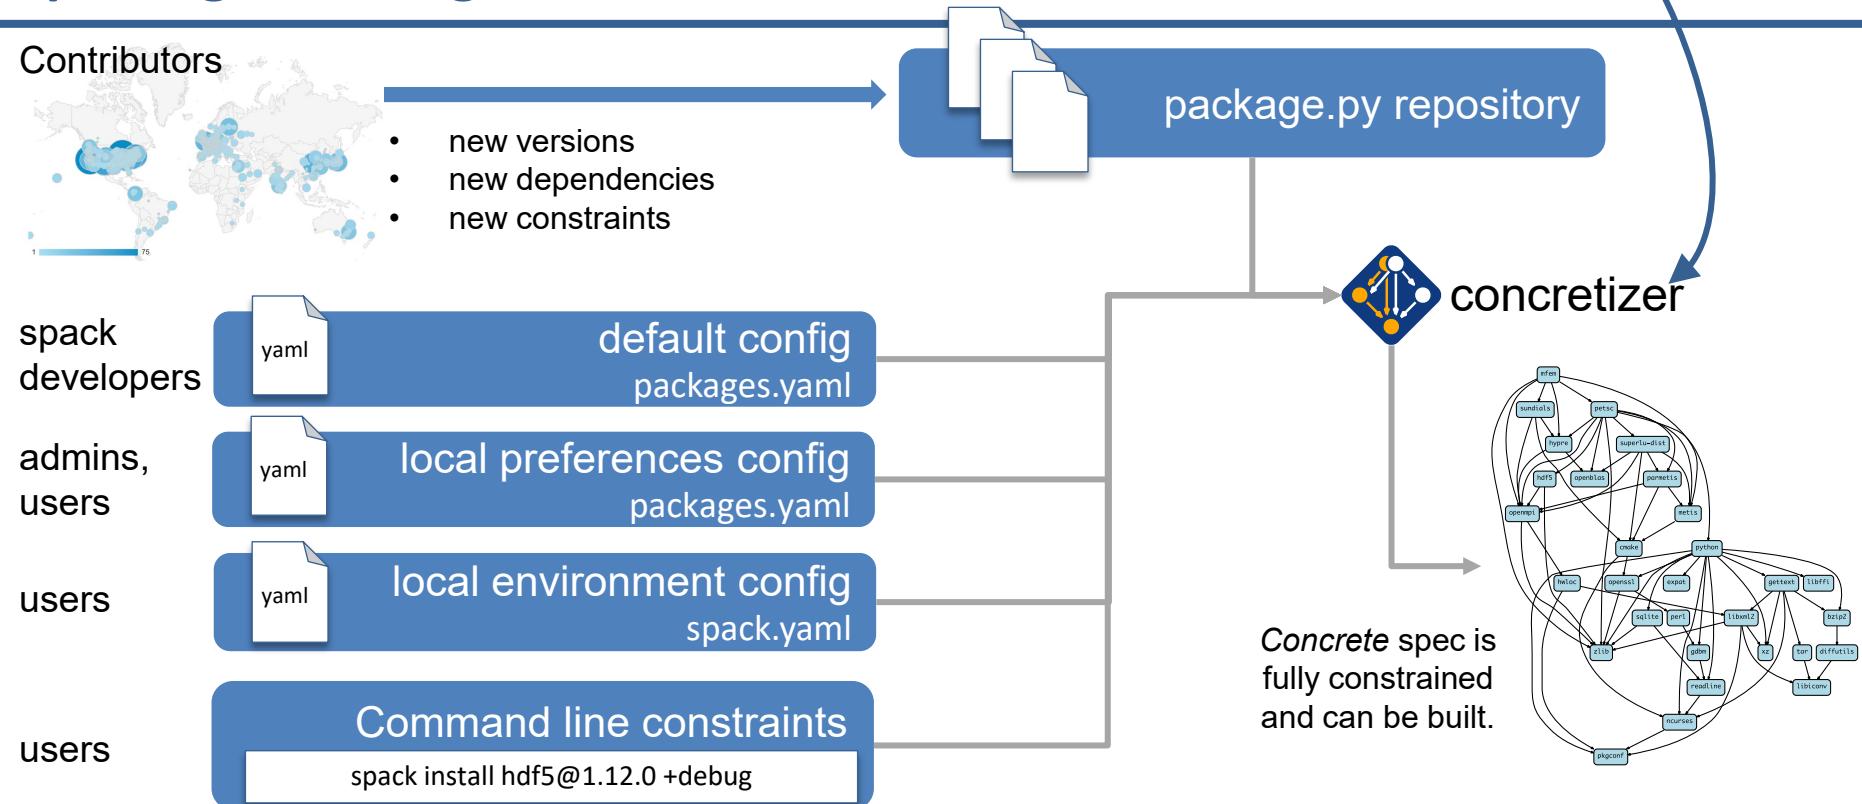
# Solving and Complexity

THE **LINUX** FOUNDATION



# The concretizer includes information from packages, configuration, and CLI

# Dependency solving is NP-hard



# Spack packages use a *lot* of (declarative) conditional logic

## CudaPackage: a mix-in for packages that use CUDA

```
class CudaPackage(PackageBase):
    variant('cuda', default=False,
            description='Build with CUDA')

    variant('cuda_arch',
            description='CUDA architecture',
            values=any_combination_of(cuda_arch_values),
            when='+cuda')

    depends_on('cuda', when='+cuda')

    depends_on('cuda@9.0:',      when='cuda_arch=70')
    depends_on('cuda@9.0:',      when='cuda_arch=72')
    depends_on('cuda@10.0:',     when='cuda_arch=75')

    conflicts('%gcc@9:', when='+cuda ^cuda@:10.2.89 target=x86_64:')
    conflicts('%gcc@9:', when='+cuda ^cuda@:10.1.243 target=ppc64le:')
```

cuda is a variant (build option)

cuda\_arch is only present  
if cuda is enabled

dependency on cuda, but only  
if cuda is enabled

constraints on cuda version

compiler support for x86\_64  
and ppc64le

There is a lot of expressive power in the Spack package DSL.

# We reimplemented Spack's concretizer using Answer Set Programming

- Was originally a greedy, custom Python algorithm
- Answer Set Programming is a *declarative* programming paradigm
  - Looks like Prolog
  - Built around modern CDCL SAT solver techniques
  - Designed for combinatorial search problems
- ASP program has 2 parts:
  1. Large list of facts generated from package recipes (problem instance)
    - 60k+ facts is typical – includes dependencies, options, etc.
  2. Small logic program (~700 lines of ASP code)
- Algorithm (the part we write) is conceptually simpler:
  - Generate facts for all possible dependencies
  - Send facts and our logic program to the solver
  - Rebuild a DAG from the results
- We're using **Clingo**, the Potassco grounder/solver package

```
%-----  
% Package: ucx  
%-----  
version_declared("uctx", "1.6.1", 0).  
version_declared("uctx", "1.6.0", 1).  
version_declared("uctx", "1.5.2", 2).  
version_declared("uctx", "1.5.1", 3).  
version_declared("uctx", "1.5.0", 4).  
version_declared("uctx", "1.4.0", 5).  
version_declared("uctx", "1.3.1", 6).  
version_declared("uctx", "1.3.0", 7).  
version_declared("uctx", "1.2.2", 8).  
version_declared("uctx", "1.2.1", 9).  
version_declared("uctx", "1.2.0", 10).  
  
variant("uctx", "thread_multiple").  
variant_single_value("uctx", "thread_multiple").  
variant_default_value("uctx", "thread_multiple", "False").  
variant_possible_value("uctx", "thread_multiple", "False").  
variant_possible_value("uctx", "thread_multiple", "True").  
  
declared_dependency("uctx", "numactl", "build").  
declared_dependency("uctx", "numactl", "link").  
node("numactl") :- depends_on("uctx", "numactl"), node("uctx").  
  
declared_dependency("uctx", "rdma-core", "build").  
declared_dependency("uctx", "rdma-core", "link").  
node("rdma-core") :- depends_on("uctx", "rdma-core"), node("uctx").  
  
%-----  
% Package: util-linux  
%-----  
version_declared("util-linux", "2.29.2", 0).  
version_declared("util-linux", "2.29.1", 1).  
version_declared("util-linux", "2.25", 2).  
  
variant("util-linux", "libuuid").  
variant_single_value("util-linux", "libuuid").  
variant_default_value("util-linux", "libuuid", "True").  
variant_possible_value("util-linux", "libuuid", "False").  
variant_possible_value("util-linux", "libuuid", "True").  
  
declared_dependency("util-linux", "pkgconfig", "build").  
declared_dependency("util-linux", "pkgconfig", "link").  
node("pkgconfig") :- depends_on("util-linux", "pkgconfig"), node("util-linux").  
  
declared_dependency("util-linux", "python", "build").  
declared_dependency("util-linux", "python", "link").  
node("python") :- depends_on("util-linux", "python"), node("util-linux").
```

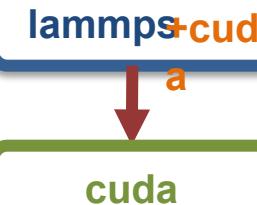
Some facts for HDF5 package

# Spack's concretizer is implemented using Answer Set Programming (ASP)

ASP looks like Prolog but is converted to SAT with optimization

Facts describe the graph

```
node("lammgs").  
node("cuda").  
variant_value("lammgs", "cuda", "True").  
depends_on("lammgs", "cuda").
```



First-order rules (with variables) describe how to resolve nodes and metadata

```
node(Dependency) :- node(Package), depends_on(Package, Dependency).
```

node("mpi")



```
node("hdf5").  
depends_on("hdf5", "mpi").
```

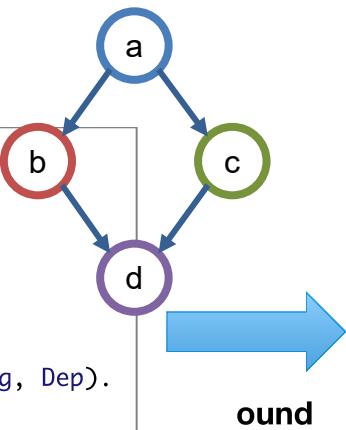
Ground  
Rule

# Grounding converts a first-order logic program into a propositional logic program, which can be solved.

```
depends_on(a, b).  
depends_on(a, c).  
depends_on(b, d).  
depends_on(c, d).
```

```
node(Dep)  
:- node(Pkg),  
    depends_on(Pkg, Dep).
```

```
% at least one is true  
1 { node(a); node(b) }.
```



ound

```
depends_on(a, b).  
depends_on(a, c).  
depends_on(b, d).  
depends_on(c, d).
```

```
node(b) :- node(a).  
node(c) :- node(a).  
node(d) :- node(c).  
node(d) :- node(b).
```

```
% at least one is true  
1 { node(a); node(b) }.
```

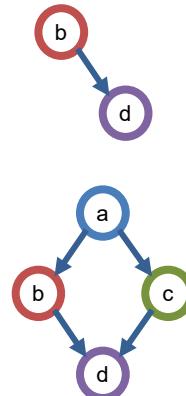


First-order Logic Program

Propositional Program

Answer 1:  
node(b)  
node(d)

Answer 2:  
node(a)  
node(b)  
node(c)  
node(d)



Stable Models (Answer Sets)

We use the *Clingo* solver from  
[potassco.org](http://potassco.org)

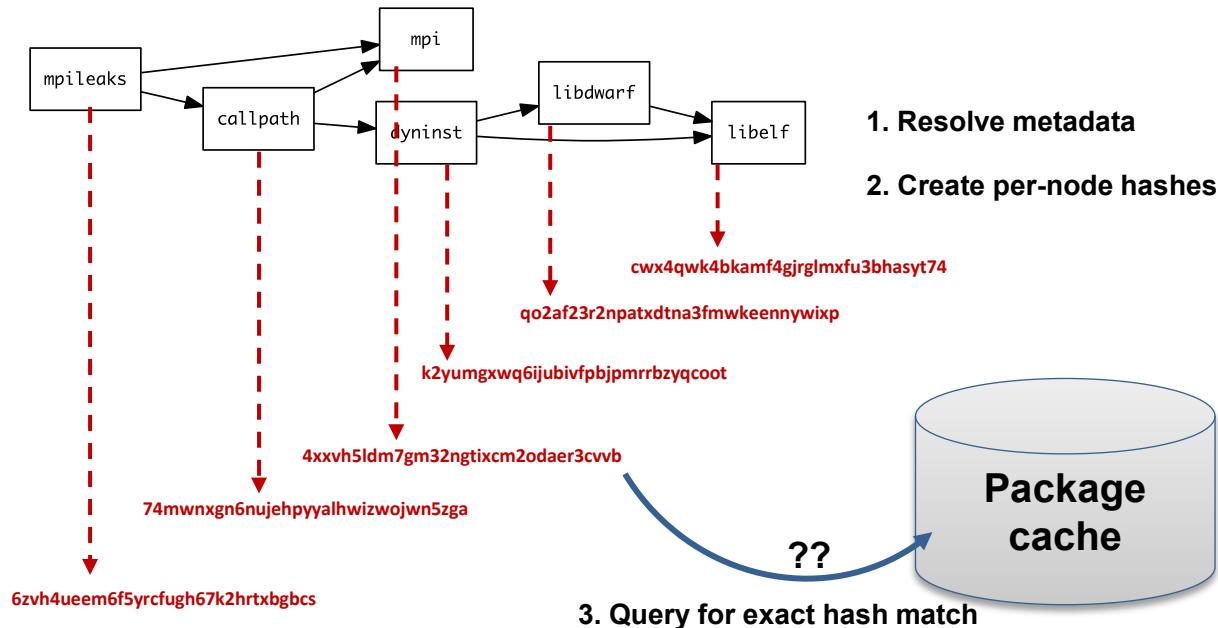
Answer 1: Only node(b) is true  
Answer 2: Both node(a) and node(b) are true

# ASP searches for *stable models* of the input program

---

- Stable models are also called *answer sets*
- A *stable model* (loosely) is a set of true atoms that can be deduced from the inputs, where every rule is idempotent.
  - Similar to fixpoints
  - Put more simply: a *set of atoms where all your rules are true!*
- Unlike Prolog:
  - Stable models contain everything that can be derived (vs. just querying values)
  - Good ways to do optimization to select the “best” stable model
  - ASP is guaranteed to complete!

# Originally, Spack only reuse builds if hashes matched



- This is now:  
spack install --fresh
- Hash matches are very sensitive to small changes
- In many cases, a satisfying cached or already installed spec can be missed
- Nix, Spack, Guix, Conan, and others reuse this way

## --reuse (now the default) was enabled by ASP

- --reuse tells the solver about all the installed packages!
- Add constraints for all installed packages, with their hash as the associated ID:

```
installed_hash("openssl", "lwatuuysmwkhahrnrywvn77icdhs6mn").  
imposed_constraint("lwatuuysmwkhahrnrywvn77icdhs6mn", "node", "openssl").  
imposed_constraint("lwatuuysmwkhahrnrywvn77icdhs6mn", "version", "openssl", "1.1.1g").  
imposed_constraint("lwatuuysmwkhahrnrywvn77icdhs6mn", "node_platform_set", "openssl", "darwin").  
imposed_constraint("lwatuuysmwkhahrnrywvn77icdhs6mn", "node_os_set", "openssl", "catalina").  
imposed_constraint("lwatuuysmwkhahrnrywvn77icdhs6mn", "node_target_set", "openssl", "x86_64").  
imposed_constraint("lwatuuysmwkhahrnrywvn77icdhs6mn", "variant_set", "openssl", "systemcerts", "True").  
imposed_constraint("lwatuuysmwkhahrnrywvn77icdhs6mn", "node_compiler_set", "openssl", "apple-clang").  
imposed_constraint("lwatuuysmwkhahrnrywvn77icdhs6mn", "node_compiler_version_set", "openssl", "apple-clang", "12.0.0").  
imposed_constraint("lwatuuysmwkhahrnrywvn77icdhs6mn", "concrete", "openssl").  
imposed_constraint("lwatuuysmwkhahrnrywvn77icdhs6mn", "depends_on", "openssl", "zlib", "build").  
imposed_constraint("lwatuuysmwkhahrnrywvn77icdhs6mn", "depends_on", "openssl", "zlib", "link").  
imposed_constraint("lwatuuysmwkhahrnrywvn77icdhs6mn", "hash", "zlib", "x2anksgssxsxa7pcnhzg5k3dhgacglze").
```



# Telling the solver to minimize builds is surprisingly simple in ASP

1. Allow the solver to *choose* a hash for any package:

```
{ hash(Package, Hash) : installed_hash(Package, Hash) } 1 :- node(Package).
```

2. Choosing a hash means we impose its constraints:

```
impose(Hash) :- hash(Package, Hash).
```

3. Define a build as something *without* a hash:

```
build(Package) :- not hash(Package, _), node(Package).
```

4. Minimize builds!

```
#minimize { 1@100, Package : build(Package) }.
```



# With and without --reuse optimization

```
(spackle):solver> spack solve -Il hdf5
=> Best of 9 considered solutions.
=> Optimization Criteria:
  Priority Criterion           Installed  ToBuild
  1   number of packages to build (vs. reuse)      -    20
  2   deprecated versions used                   0    0
  3   version weight                           0    0
  4   number of non-default variants (roots)     0    0
  5   preferred providers for roots            0    0
  6   default values of variants not being used (roots) 0    0
  7   number of non-default variants (non-roots) 0    0
  8   preferred providers (non-roots)          0    0
  9   compiler mismatches                     0    0
 10  OS mismatches                          0    0
 11  non-preferred OS's                      0    0
 12  version badness                        0    2
 13  default values of variants not being used (non-roots) 0    0
 14  non-preferred compilers                 0    0
 15  target mismatches                     0    0
 16  non-preferred targets                  0    0

- zznqfs3 hdf5@1.10.7%apple-clang@13.0.0~cxx~fortran~hl~ipo~java+mpi+shared+szip~threadsafe+tools api=default b
  ^cmake@3.21.4%apple-clang@13.0.0~doc+ncurses+openssl+owlbins+qt build_type=Release arch=darwin-bigsur-skylake
  ^ncurses@6.2%apple-clang@13.0~symlinks+termlib abi=none arch=darwin-bigsur-skylake
  ^pkgconf@1.8.0%apple-clang@13.0.0 arch=darwin-bigsur-skylake
  ^openssl@1.1.1%apple-clang@13.0.0~docs certsys system arch=darwin-bigsur-skylake
  ^perl@5.34.0%apple-clang@13.0.0~cpam+shared+threads arch=darwin-bigsur-skylake
  ^berkeley-db@18.1.40%apple-clang@13.0.0~cxx+docs+stl patches=b231fcc4d5cff05e5c3a4814f
  ^bzzip2@0.8%apple-clang@13.0.0~debug+pic+shared arch=darwin-bigsur-skylake
  ^diffutils@3.8%apple-clang@13.0.0 arch=darwin-bigsur-skylake
  ^libiconv@1.16%apple-clang@13.0.0~libs=shared,static arch=darwin-bigsur-skylake
  ^gdbm@1.19%apple-clang@13.0.0 arch=darwin-bigsur-skylake
  ^readline@8.1%apple-clang@13.0.0 arch=darwin-bigsur-skylake
  ^zlib@1.2.11%apple-clang@13.0.0~optimize+pic+shared arch=darwin-bigsur-skylake
  ^xz@5.2.5%apple-clang@13.0.0~pic libs=shared,static arch=darwin-bigsur-skylake
  ^openmp@4.1.1%apple-clang@13.0.0~atomics+cuda-cxx-cxx_exceptions+gfps+internal-hwloc+java+legacy
  ^hwloc@2.6.0%apple-clang@13.0.0~cairo+cuda+gl+libudev+libxml2+netloc+nvml+opencl+pci+rocm+sh
  ^libxml2@2.9.12%apple-clang@13.0.0~python arch=darwin-bigsur-skylake
  ^xz@5.2.5%apple-clang@13.0.0~pic libs=shared,static arch=darwin-bigsur-skylake
  ^libevent@2.1.12%apple-clang@13.0.0~openssl arch=darwin-bigsur-skylake
  ^openssl@8.7p1%apple-clang@13.0.0 arch=darwin-bigsur-skylake
  ^libedit@0.1-20210216%apple-clang@13.0.0 arch=darwin-bigsur-skylake
```

Pure hash-based reuse: all misses

```
(spackle):spack> spack solve --reuse -Il hdf5
=> Best of 10 considered solutions.
=> Optimization Criteria:
  Priority Criterion           Installed  ToBuild
  1   number of packages to build (vs. reuse)      -    4
  2   deprecated versions used                   0    0
  3   version weight                           0    0
  4   number of non-default variants (roots)     0    0
  5   preferred providers for roots            0    0
  6   default values of variants not being used (roots) 0    0
  7   number of non-default variants (non-roots) 2    0
  8   preferred providers (non-roots)          0    0
  9   compiler mismatches                     0    0
 10  OS mismatches                          0    0
 11  non-preferred OS's                      0    0
 12  version badness                        6    0
 13  default values of variants not being used (non-roots) 1    0
 14  non-preferred compilers                 15   4
 15  target mismatches                     0    0
 16  non-preferred targets                  0    0

- yfkfnsp hdf5@1.10.7%apple-clang@12.0.5~cxx~fortran~hl~ipo~java+mpi+shared+szip~threadsafe+tools api=default b
  ^cmake@21.1%apple-clang@12.0.5~doc+ncurses+openssl+owlbins+qt build_type=Release arch=darwin-bigsur-skylake
  ^ncurses@6.2%apple-clang@12.0.5~symlinks+termlib abi=none arch=darwin-bigsur-skylake
  ^openssl@1.1.1%apple-clang@12.0.5~docs+systemcerts arch=darwin-bigsur-skylake
  ^perl@2.1.11%apple-clang@12.0.5~optimize+pic+shared arch=darwin-bigsur-skylake
  ^openmp@4.1.1%apple-clang@12.0.5~atomics+cuda-cxx-cxx_exceptions+gfps+internal-hwloc+java+leg
  ^hwloc@2.6.0%apple-clang@12.0.5~cairo+cuda+gl+libudev+libxml2+netloc+nvml+opencl+pci+rocm+
  ^libxml2@2.9.12%apple-clang@12.0.5~python arch=darwin-bigsur-skylake
  ^xz@5.2.5%apple-clang@12.0.5~pic libs=shared,static arch=darwin-bigsur-skylake
  ^libiconv@1.16%apple-clang@12.0.5~libs=shared,static arch=darwin-bigsur-skylake
  ^gdbm@1.19%apple-clang@12.0.5~pic libs=shared,static arch=darwin-bigsur-skylake
  ^readline@8.1%apple-clang@12.0.5~pic libs=shared,static arch=darwin-bigsur-skylake
  ^pkgconf@1.8.0%apple-clang@12.0.5 arch=darwin-bigsur-skylake
  ^libevent@2.1.12%apple-clang@12.0.5~openssl arch=darwin-bigsur-skylake
  ^libedit@0.1-20210216%apple-clang@12.0.5 arch=darwin-bigsur-skylake
  ^perl@5.34.0%apple-clang@12.0.5~cpam+shared+threads arch=darwin-bigsur-skylake
  ^berkeley-db@18.1.40%apple-clang@12.0.5~cxx+docs+stl patches=b231fcc4d5cff05e5c3a4814f
  ^bzzip2@1.0.8%apple-clang@12.0.5~debug+pic+shared arch=darwin-bigsur-skylake
  ^gdbm@1.19%apple-clang@12.0.5 arch=darwin-bigsur-skylake
  ^readline@8.1%apple-clang@12.0.5 arch=darwin-bigsur-skylake
```

With reuse: 16 packages were reusable



# Recent Challenges

THE **LINUX** FOUNDATION

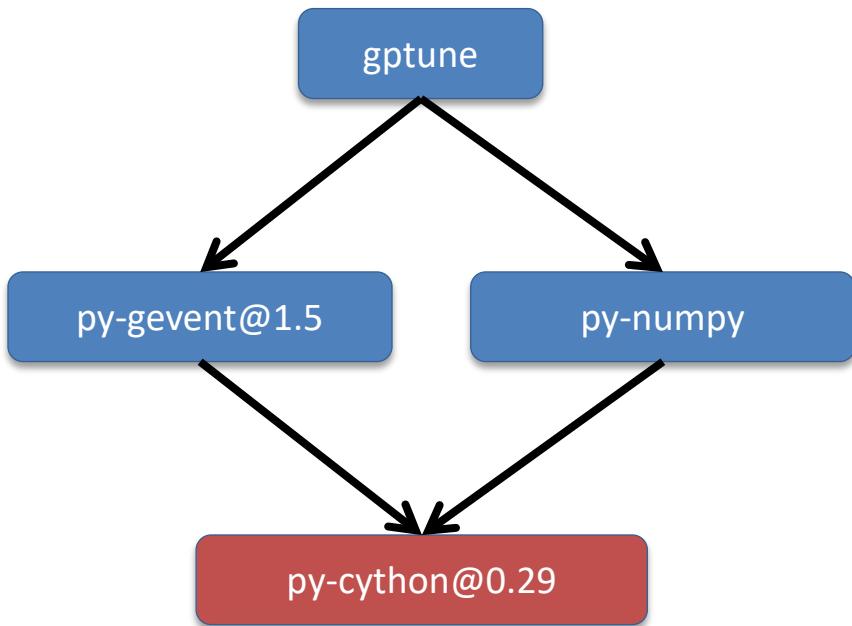


# Some stats on problem sizes

- Main logic program is:
  - ~250 rules
  - 20 optimization criteria
  - 933 lines of ASP code
- Problem instances can vary quite a bit
  - Common dependencies get us some magic numbers
  - gmake's optional dependency on guile makes most solves consider at least 527 packages
  - gnuconfig is notably very simple ☺

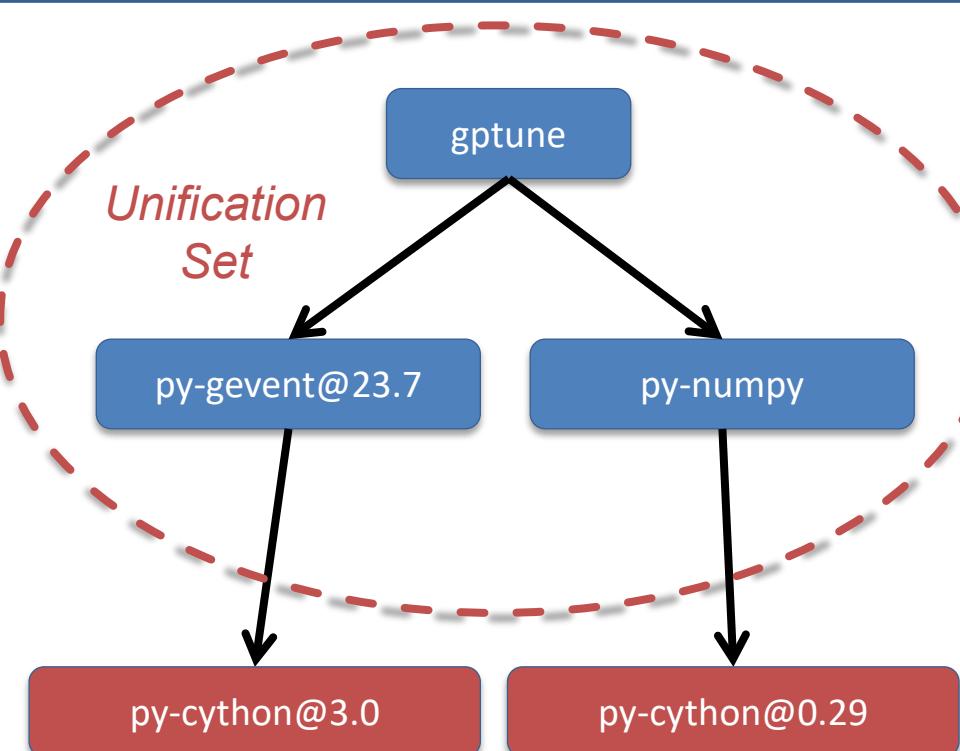
Package	Possible dependencies	Facts
gnuconfig	1	150
zlib	527	30,095
gmake	527	30,160
openmpi	527	109,021
qt	527	109,029
trilinos	694	224,142
root	699	146,372
mfem	714	273,078
r-condop	774	142,212
warpx	819	319,374
exawind	820	322,535

# We have been working to make our solver more flexible



- Only one configuration per package allowed in the DAG
- Ensures ABI compatibility but is too restrictive
- In the example py-numpy needs to use py-cython@0.29 as a build tool
- That enforces using an old py-gevent, because newer versions depend on py-cython@3.0 or greater

# Objective: dependency splitting



- The constraint on build dependencies can be relaxed, without compromising the ABI compatibility
- Having a single configuration of a package is now enforced on unification sets
- These are the set of nodes used together at runtime (the one shown is for gptune)
- This allows us to use the latest version of py-gevent, because now we can have two versions of py-cython

# We want to dynamically “split” nodes when needed

Start with deducing single dependency nodes:

```
node(DependencyName)
:- dependency_holds(PkgName, DependencyName)
```

Want to allow solver to **choose** to duplicate a node: Converted node identifier  
from **name** to (**name, id**)

```
1 {
  depends_on(PkgNode, node(0..Y-1, DepNode), Type)
  : max_dupes(DepNode, Y)
} 1
:- dependency_holds(PkgNode, DepNode).
```



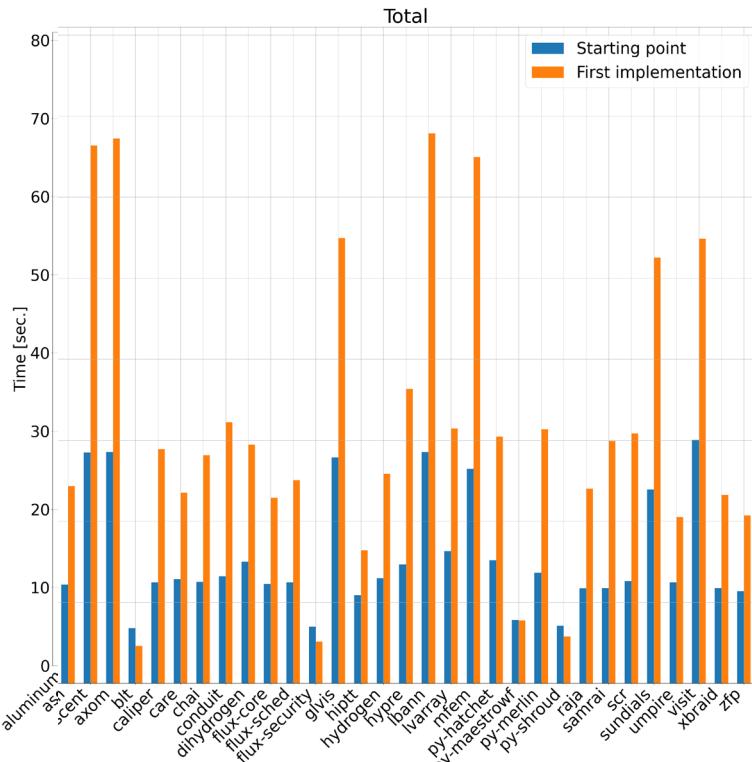
# Generic package metadata can be used with any duplicate node

```
pkg_fact("alsa-lib",version_declared("1.2.3.2",0,"package_py")).  
pkg_fact("alsa-lib",version_declared("1.2.2",1,"package_py")).  
pkg_fact("alsa-lib",version_declared("1.1.4.1",2,"package_py")).
```

```
pkg_fact("alsa-lib",condition(20)).  
condition_reason(20,"alsa-lib depends on python when +python").  
pkg_fact("alsa-lib",condition_trigger(20,15)).
```

- Facts that come from package descriptions can be used with all duplicate nodes
- We now have to ground multiple copies of most of our rules
- Performance still scales with total number of possible nodes
  - Small numbers of duplicates don't really explode the solve

# First try at allowing duplicates in a single solve



Increased solve times by  
=> 2x in some cases

# It turns out that cycle detection in the solver is *expensive*

```
path(A, B) :- depends_on(A, B).  
path(A, C) :- path(A, B), depends_on(B, C).
```

% this constraint says "no cycles"

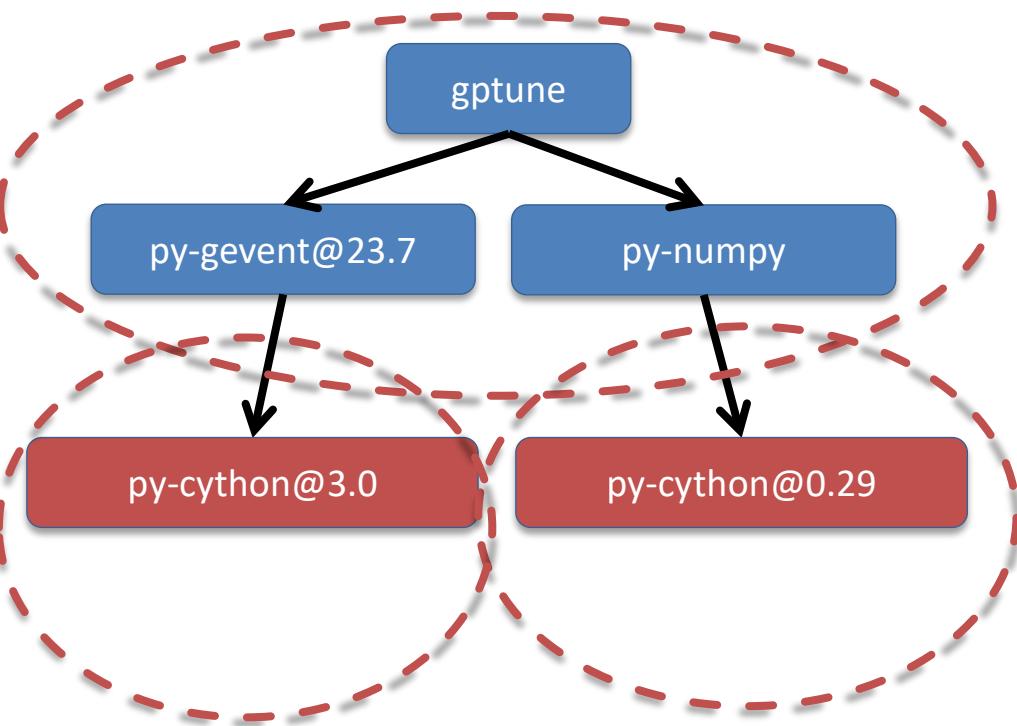
```
:- path(A, B), path(B, A).
```

- Has to maintain path() predicate representing paths between nodes
- Cycles are actually rare in solutions
  - Switched to post-processing for cycle detection
  - Only do expensive solve if a cycle is detected in a solution
- Similar issue arose for variant propagation in graph
  - Fixed by reworking variant propagation not to track paths

50%+ improvement  
in solve time

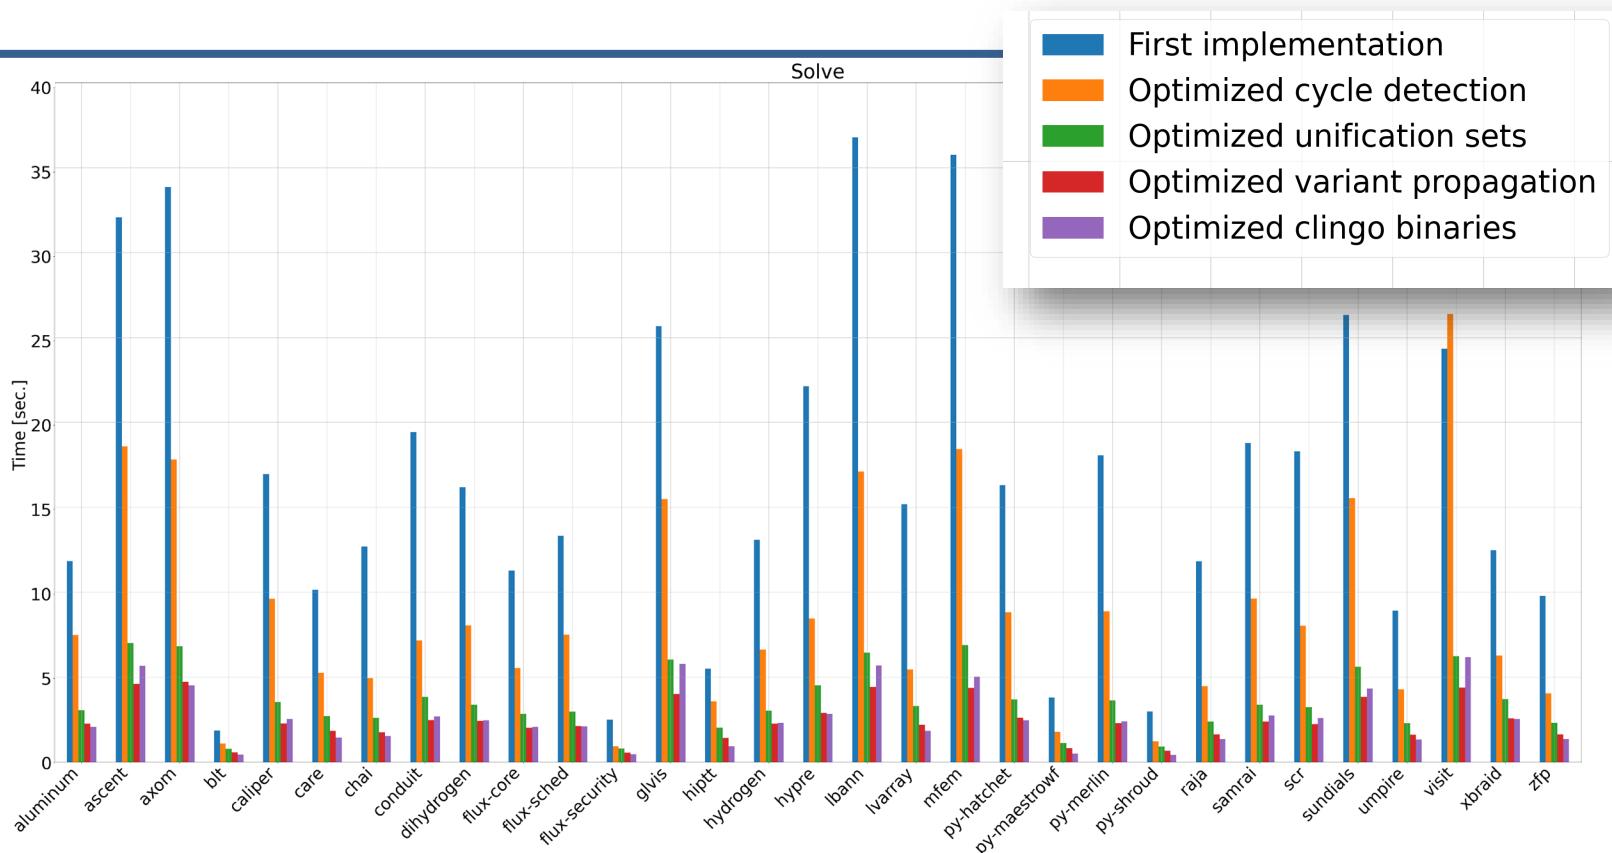


# Unification sets can be expensive too

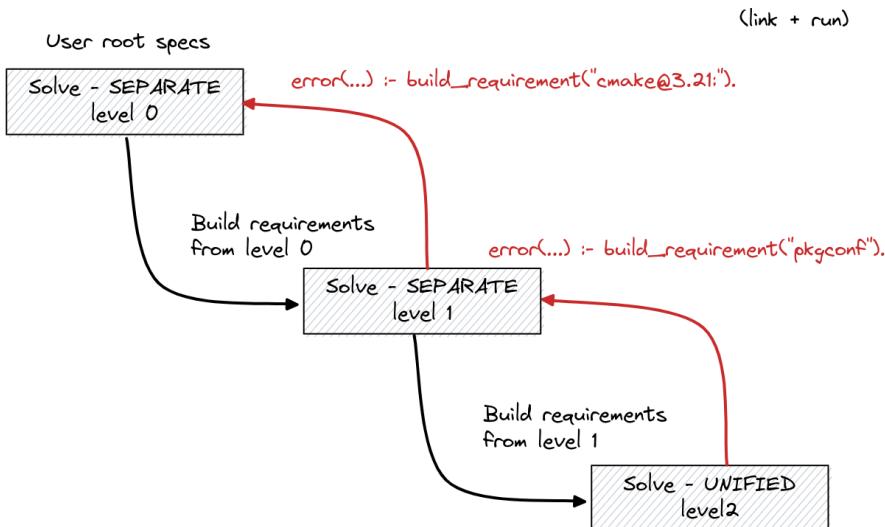


- Unification set creation was originally recursive for *any* build dependencies
  - Ends up blowing up grounding
- Mitigation:
  - Only create new sets for explicitly *marked* build tools
  - Transitive build dependencies that are not from marked build tools go into a *common* unification set
- Need better heuristics to split when necessary for full generality

# Optimizations: Solve Time



# It was not trivial to come up with this model



- In addition to this “coupled” method, we tried an iterative version with multiple solves
- Multiple solves had some disadvantages:
  - Slower due to overhead of multiple solves
  - Not coupled, so feedback from solve to solve was awkward
  - Packagers needed to “help” the solver
- Requiring packagers to provide solve hints in packages isn’t practical

# Recent Work

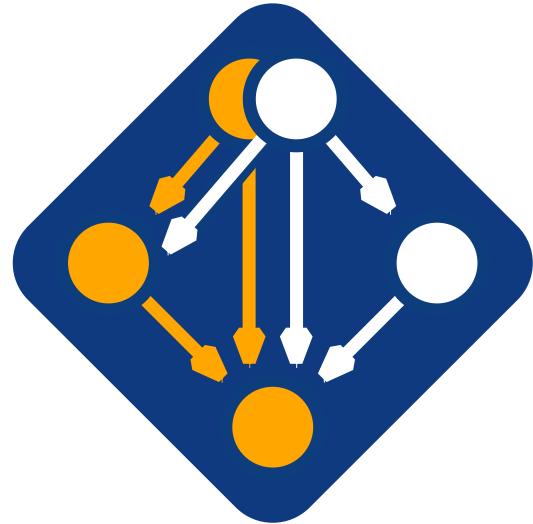
THE **LINUX** FOUNDATION



# Spack v0.22.0 is out!

## Highlights:

1. Compiler runtime dependencies
  - gcc-runtime, intel-oneapi-runtime, libgfortran, libc
  - OS compatibility on linux now uses libc version, not OS tag
2. Improved spack find UI for Environments
3. Improved command-line string quoting
4. Revert default spack install behavior to --reuse
5. More control over reused specs
6. New redistribute() directive
7. New conflict: and prefer: syntax for package preferences
8. include\_concrete: in environments
9. python-venv isolation



[github.com/spack/spack](https://github.com/spack/spack)

## Full release notes:

<https://github.com/spack/spack/releases/tag/v0.22.0>

# We've made a lot of progress on compiler dependencies

---

- Compiler *runtime libraries* represented in the graph
  - C++, Fortran runtimes
- libc is now represented in dependency graphs on Linux
  - No more need to rely on OS tag for compatibility information
- Reuse binaries *without* their compiler needing to be configured locally
- Improved buildcache hit rate using libraries for compatibility

# Compilers can now model their own runtimes

- New method, runtime\_constraints, for injecting runtimes into graphs
- Currently supported for gcc, intel-oneapi
  - still working on others
- Allows solver to take libstdc++, fortran runtime compatibility into account.
- Example:
  - Intel compilers now (correctly) depend on gcc-runtime

```
class Gcc(AutotoolsPackage, GNUMirrorPackage):
    # ...

    @classmethod
    def runtime_constraints(cls, *, spec, pkg):
        """Callback function to inject runtime-related rules into the solver.

        Rule-injection is obtained through method calls of the ``pkg`` argument.

        Documentation for this function is temporary. When the API will be in its final state,
        we'll document the behavior at https://spack.readthedocs.io/en/latest/

        Args:
            spec: spec that will inject runtime dependencies
            pkg: object used to forward information to the solver
        """
        pkg("*").depends_on(
            "gcc-runtime",
            when="%gcc",
            type="link",
            description="If any package uses %gcc, it depends on gcc-runtime",
        )
        pkg("*").depends_on(
            f"gcc-runtime@{str(spec.version)}:",
            when=f"%{str(spec)}",
            type="link",
            description=f"If any package uses %{str(spec)}, "
            f"it depends on gcc-runtime@{str(spec.version)}:",
        )
```

# Packages now declare the languages they depend on

- Languages are *almost* virtuals
  - HDF5 package depends on `cxx` and `fortran`
  - Handled specially internally, until compilers are nodes
- Imply a compiler *and* compiler package can specify runtime libraries to inject
- Allows solver to mix compilers *correctly*
  - Runtimes are unified like other nodes
  - *Package authors* can model toolchain properties
  - probably not for most package authors, but very powerful
- TBD:
  - Other runtimes like clang libraries and OpenMP
  - Compilers as nodes in the graph

```
class Hdf5(CMakePackage):  
    """HDF5 is a data model, library, and file format for storing and managing  
    data. It supports an unlimited variety of datatypes, and is designed for  
    flexible and efficient I/O and for high volume and complex data.  
    """  
  
    homepage = "https://portal.hdfgroup.org"  
    url = "https://support.hdfgroup.org/ftp/HDF/releases/hdf5-1.14/hdf5-1.14.4  
    list_url = "https://support.hdfgroup.org/ftp/HDF/releases"  
    list_depth = 3  
    git = "https://github.com/HDFGroup/hdf5.git"  
    maintainers("lrknox", "brtnfld", "byrnHDF", "gheber", "hyoklee", "lkurz")  
  
    tags = ["e4s", "windows"]  
    executables = ["^h5cc$", "^h5pcc$"]  
  
    test_requires_compiler = True  
  
    license("custom")  
  
    depends_on("cxx", type="build", when="+cxx")  
    depends_on("fortran", type="build", when="+fortran")
```

# We've also added libc as a dependency

- libc is a virtual
  - glibc and musl packages are providers
  - (nearly) every graph has libc in it, via the compiler
  - Can be external or built by Spack
- We are *not* building libc for every stack in Spack
  - Automatically detect system libc version
  - Add a node to the graph to be used for binary compatibility
- No longer using OS tags for buildcaches
  - Now use libc for this
  - *many* more buildcache hits

```
(py311) culpo@nivola:~/PycharmProjects/spack$ spack concretize -f --re
==> Concretized hdf5~mpi
- tnqdhsn  hdf5@1.14.3%gcc@9.4.0~cxx~fortran~hl~ipo~java~map~mpi+sh
- gdviueh   ^cmake@3.27.9%gcc@8.5.0~doc+ncurses+ownlibs build_sy
- i5cd2j j   ^curl@8.6.0%gcc@8.5.0~gssapi~ldap~libidn2~librtm
- icqajq4   ^nghttp2@1.57.0%gcc@8.5.0 build_system=autot
- xl7h3wb   ^openssl@3.2.1%gcc@8.5.0~docs+shared build_s
- rlipoky   ^ca-certificates-mozilla@2023-05-30%gcc@8.5.0
- 45hdvpf   ^perl@5.38.0%gcc@8.5.0+cpanm+opcode+open
- qvzagc5   ^berkeley-db@18.1.40%gcc@8.5.0+cxx+dl
- ioufq6d   ^bzip2@1.0.8%gcc@8.5.0~debug~pic+sha
- enaxy2l   ^diffutils@3.10%gcc@8.5.0 build_
- czvftrb   ^libiconv@1.17%gcc@8.5.0 buil
- ku6webf   ^gdbm@1.23%gcc@8.5.0 build_system=au
- 3tzxgdp   ^readline@8.2%gcc@8.5.0 build_sy
- llqwd2j   ^gcc-runtime@8.5.0%gcc@8.5.0 build_system=generi
[e] fue5ca2  ^glibc@2.28%gcc@8.5.0 build_system=autotools arc
- sxb2sl6   ^ncurses@6.4%gcc@8.5.0~svmlinks+termlib abi=none
- ucn3h     ^gcc-runtime@9.4.0%gcc@9.4.0 build_system=generic ar
[e] 37zrmp4   ^glibc@2.31%gcc@9.4.0 build_system=autotools arch=li
- vsjxwea   ^make@4.4.1%gcc@8.5.0~gui~lto build_system=generic ar
- o76bf47   ^pkgconf@1.9.5%gcc@8.5.0 build_system=autotools arch=
- jkwqkvs   ^zlib-ng@2.1.6%gcc@8.5.0+compat+new_strategies+opt+pk
```

# Libc modeling makes for a much better buildcache experience

- Currently on develop (emacs 100% from binary):

```
(py311) culpo@nivola:~/PycharmProjects/spack$ spack install emacs
[+] /usr (external glibc-2.17-2hhcy7kzv3lfqcaschwup4uysp4hoy)
[+] /home/culpo/PycharmProjects/spack/opt/spack/linux-centos7-x86_64_v3/gcc-10.2.1/gcc-runtime-10.2.1-4gmidou73wvttvhun564olcopuijl2i
[+] /home/culpo/PycharmProjects/spack/opt/spack/linux-centos7-x86_64_v3/gcc-10.2.1/gmp-6.2.1-nkhm7cmp6samsykyksuqda77xbxnibfn
[+] /home/culpo/PycharmProjects/spack/opt/spack/linux-centos7-x86_64_v3/gcc-10.2.1/berkeley-db-18.1.40-thoi4z7lozaednxhjd4ojhw2lcsgo5g
[+] /home/culpo/PycharmProjects/spack/opt/spack/linux-centos7-x86_64_v3/gcc-10.2.1/gmake-4.4.1.ucqstlhik7pmv5ijyckmr6mxv46vgeuj
=> Installing nasm-2.15.05-e2kvtpqiaava2osr3jbyptdufjgov4dgc [6/39]
=> Fetching https://binaries.spack.io/develop/developer-tools-manylinux2014/build_cache/linux-centos7-x86_64_v3-gcc-10.2.1-nasm-2.15.05-e2kvtpqiaava2osr3jbyptdufjgov4dgc.spec.json.sig
gpg: Signature made mar 23 apr 2024, 14:38:11 CEST
gpg:           using RSA key D2C7EB3F2B05FA86590D293C04001B2E3DB0C723
gpg: Good signature from "Spack Project Official Binaries <maintainers@spack.io>" [ultimate]
=> Fetching https://binaries.spack.io/develop/developer-tools-manylinux2014/build_cache/linux-centos7-x86_64_v3/gcc-10.2.1/nasm-2.15.05/linux-centos7-x86_64_v3-gcc-10.2.1-nasm-2.15.05-e2kvtpqiaava2osr3jbyptdufjgov4dgc.spec.json.sig
gpg: Signature made mar 23 apr 2024, 14:38:11 CEST
gpg:           using RSA key D2C7EB3F2B05FA86590D293C04001B2E3DB0C723
=> Extracting nasm-2.15.05-e2kvtpqiaava2osr3jbyptdufjgov4dgc from binary cache
=> nasm: Successfully installed nasm-2.15.05-e2kvtpqiaava2osr3jbyptdufjgov4dgc
  Search: 0.00s.  Fetch: 2.61s.  Install: 0.17s.  Extract: 0.11s.  Relocate: 0.04s.  Total: 2.78s
[+] /home/culpo/PycharmProjects/spack/opt/spack/linux-centos7-x86_64_v3/gcc-10.2.1/nasm-2.15.05-e2kvtpqiaava2osr3jbyptdufjgov4dgc
=> Installing pcre-8.45-xi42ks7asbq2sqmdc7bdsmhzzhlie6m4 [7/39]
=> Fetching https://binaries.spack.io/develop/developer-tools-manylinux2014/build_cache/linux-centos7-x86_64_v3-gcc-10.2.1-pcre-8.45-xi42ks7asbq2sqmdc7bdsmhzzhlie6m4.spec.json.sig
gpg: Signature made mar 23 apr 2024, 14:38:12 CEST
gpg:           using RSA key D2C7EB3F2B05FA86590D293C04001B2E3DB0C723
gpg: Good signature from "Spack Project Official Binaries <maintainers@spack.io>" [ultimate]
=> Fetching https://binaries.spack.io/develop/developer-tools-manylinux2014/build_cache/linux-centos7-x86_64_v3/gcc-10.2.1/pcre-8.45/linux-centos7-x86_64_v3-gcc-10.2.1-pcre-8.45-xi42ks7asbq2sqmdc7bdsmhzzhlie6m4
=> Extracting pcre-8.45-xi42ks7asbq2sqmdc7bdsmhzzhlie6m4 from binary cache
=> pcre: Successfully installed pcre-8.45-xi42ks7asbq2sqmdc7bdsmhzzhlie6m4
  Search: 0.00s.  Fetch: 2.34s.  Install: 0.18s.  Extract: 0.14s.  Relocate: 0.03s.  Total: 2.52s
[+] /home/culpo/PycharmProjects/spack/opt/spack/linux-centos7-x86_64_v3/gcc-10.2.1/pcre-8.45-xi42ks7asbq2sqmdc7bdsmhzzhlie6m4
=> Installing tree-sitter-0.22.2-x5wmc6tind5qxo2zrswb7vyj3xuhak2d [8/39]
=> Fetching https://binaries.spack.io/develop/developer-tools-manylinux2014/build_cache/linux-centos7-x86_64_v3-gcc-10.2.1-tree-sitter-0.22.2-x5wmc6tind5qxo2zrswb7vyj3xuhak2d.spec.json.sig
gpg: Signature made mar 23 apr 2024, 14:38:22 CEST
gpg:           using RSA key D2C7EB3F2B05FA86590D293C04001B2E3DB0C723
gpg: Good signature from "Spack Project Official Binaries <maintainers@spack.io>" [ultimate]
=> Fetching https://binaries.spack.io/develop/developer-tools-manylinux2014/build_cache/linux-centos7-x86_64_v3/gcc-10.2.1/tree-sitter-0.22.2/linux-centos7-x86_64_v3-gcc-10.2.1-tree-sitter-0.22.2-x5wmc6tind5qxo2zrswb7vyj3xuhak2d
=> Extracting tree-sitter-0.22.2-x5wmc6tind5qxo2zrswb7vyj3xuhak2d from binary cache
=> tree-sitter: Successfully installed tree-sitter-0.22.2-x5wmc6tind5qxo2zrswb7vyj3xuhak2d
  Search: 0.00s.  Fetch: 2.24s.  Install: 0.00s.  Extract: 0.03s.  Relocate: 0.04s.  Total: 2.23s
```



# Roadmap

THE **LINUX** FOUNDATION



# Roadmap for this year

---

1. Finish compiler dependencies
  - Compilers will appear as build dependencies
  - All special compiler logic generified for any build dependency
  - Enable bootstrapping compilers from buildcache
  - Continue to flesh out low-level compiler runtime libraries
2. Harden public build caches and make them more broadly compatible
3. Enable bare-metal MPI/CUDA/ROCm installations
4. Make the build cache on by default; build only what we need from source
5. Automatic Python package generation
6. Continue to improve Windows support
7. Speed improvements for concretization and metadata management

**Lots of exciting work ahead!**

# Autogenerating Python packages

- Python ecosystem is a lot to maintain
  - Increasingly, we need to update faster than we can sustain through pull requests
  - Reviewing version ranges has been painful
- We added features to our version system to support Python better
  - Prerelease, alpha, beta version comparison
- We can now generate Python packages from PyPI metadata
  - Only for pure python packages
  - Native packages will be maintained as regular Spack packages
  - Overlay on top of auto-generated python packages

```
from spack.package import *

class PyBlack(PythonPackage):
    version("24.3.0", sha256="a0c9c4a0771afc6919578cec71ce82a3e31e054904e7197deacbc9382671c"
version("24.2.0", sha256="bce4f26c27c3435e4dace4815cbcb008b87e167e3bf4ee47ccdc5ce906eb4c"
version("24.1.1", sha256="48b5760dcbfe5cf97fd4ba2394681f3a81514c6abbba45b50da7ac8fbcb6c"
version("24.1.0", sha256="30fbf768cd4f4576598b1db020243fafe9a227ef880d1a12230c643cefef"
version("23.12.1", sha256="4ce3e14eb809509188014d9ea1c456a910d5b5cbf43a09ffef7e024b3c"
version("23.12.0", sha256="330a327b422aca0634ecd15985c1c7fd7bdb5b5a2ef8aa9888a82ebe9"
version("23.11.0", sha256="4c68855825ff432d19729846f971bc4d666ce98492e502013bcaca4d9"
version("23.10.1", sha256="1f8ce31675342ff68749c65a5f7844631a18c8679dfdf3ca9dc1a289979c"
version("23.10.0", sha256="31bf9f87b277a6800e99d2985eda08807c087973eaa699da5f0c6def6b7c"
version("23.9.1", sha256="24bb603ff5c6d9ea08a8886f6977eae858e1f340d7260cf56d70a49823236bc

variant("colorama", default=False)
variant("d", default=False)
variant("jupyter", default=False)
variant("uvloop", default=False)

with default_args(type="run"):
    depends_on("py-aiohttp@3.7.4:", when="@23.12:23.12.0,24:24.1.0 platform=linux")
    depends_on("py-aiohttp@3.7.4:", when="@23.12:23.12.0,24:24.1.0 platform=freebsd")
    depends_on("py-aiohttp@3.7.4:", when="@23.12:23.12.0,24:24.1.0 platform=darwin")
    depends_on("py-aiohttp@3.7.4:", when="@23.12:23.12.0,24:24.1.0 platform=cray")
    depends_on("py-aiohttp@3.7.4:", when="@21.10-beta0:21,22.10:+d")
    depends_on("py-click08.0.0:", when="@22.10:")
    depends_on("py-colorama@0.4.3:", when="@20:21,22.10:+colorama")
    depends_on("py-ipython@7.8:", when="@21.8-beta0:21,22.10:+jupyter")
    depends_on("py-mypy-extensions@0.4.3:", when="@20:21,22.10:")
    depends_on("py-packaging@22:", when="@23.1.0:")
    depends_on("py-pathspec@0.9.+", when="@22.10:")
    depends_on("py-platformdirs@2.0.0:", when="@21.8-beta0:21,22.10:")
    depends_on("py-tokenize@0.3.2:", when="@21.8-beta0:21,22.10:+jupyter")
    depends_on("py-tomli@1.1:", when="@22.10: python@3.10")
    depends_on("py-typing-extensions@0.1.1:", when="@23.9: ^python@3.10")
    depends_on("py-uvloop@0.15.2:", when="@21.5-beta2:21,22.10:+uvloop")
```

# Generating Python

```
$ ./src/package.py generate --clean spack_requirements.txt
```

- Had to rework version system to handle some frequent python conventions
  - Prereleases
  - Alpha, beta
  - Release candidate
- Now PyPI metadata can map to Spack depends\_on() constraints
- Some caveats
  - Only doing this for pure python packages
  - Wheels (and PyPI database) lack build dependency info
  - May not be able to use spack develop with these

```
from spack.package import *

class PyBlack(PythonPackage):
    version("24.3.0", sha256="a0c9c4a0771afc6919578cec71ce82a3e31e054904e7197deacbc9382671c"
version("24.2.0", sha256="bce4f25c27c3435e4dace4815bcb2008b87e167e3bf4ee47ccdc5ce906eb4"
version("24.1.1", sha256="48b5760dcbe5cf97fd4fba23946681f3a81514c6ab8a45b50da67ac8fbcc6c"
version("24.1.0", sha256="30fbf768cd4f4576598b1db0202413fafaea9a227ef808d1a12230c643cef5"
version("23.12.1", sha256="4ce3ef14ebe8d9509188014d96af1c456a910d5b5cbf434a09fef7e024b3c"
version("23.12.0", sha256="330a327b422aca0634ecd115985c1c7fd7bdb5b5a2ef8aa9888a82e2bebe9"
version("23.11.0", sha256="4c68855825ff432d197229846f971bc4d6666ce90492e5b02013bcaca4d9"
version("23.10.1", sha256="1f8ce316753428ff68749c65a5f7844631aa18c8679df3ca9dc1a289979c"
version("23.10.0", sha256="31b9f87b277a68d0e99d2905edae08807c007973aa609da5f0c62def6b7c"
version("23.9.1", sha256="24b6b3ff5c6d9ea08a8888f6977eae858e1f340d7260cf56d70a49823236b7

variant("colorama", default=False)
variant("d", default=False)
variant("jupyter", default=False)
variant("uvloop", default=False)

with default_args(type="run"):
    depends_on("py-aiohttp@3.7.4:", when="@23.12:23.12.0,24:24.1.0 platform=linux")
    depends_on("py-aiohttp@3.7.4:", when="@23.12:23.12.0,24:24.1.0 platform=freebsd")
    depends_on("py-aiohttp@3.7.4:", when="@23.12:23.12.0,24:24.1.0 platform=darwin")
    depends_on("py-aiohttp@3.7.4:", when="@23.12:23.12.0,24:24.1.0 platform=cray")
    depends_on("py-aiohttp@3.7.4:", when="@21.10-beta0:21,22.10:+d")
    depends_on("py-click@0.8.0:", when="@22.10:")
    depends_on("py-colorama@0.4.3:", when="@20:21,22.10:+colorama")
    depends_on("py-ipython@0.7.8:", when="@21.8-beta0:21,22.10:+jupyter")
    depends_on("py-mypy-extensions@0.4.3:", when="@20:21,22.10:")
    depends_on("py-packaging@22:", when="@23.1.0:")
    depends_on("py-pathspec@0.9:", when="@22.10:")
    depends_on("py-platformdirs@2.0.0:", when="@21.8-beta0:21,22.10:")
    depends_on("py-tokenize-rt@3.2:", when="@21.8-beta0:21,22.10:+jupyter")
    depends_on("py-tomli@1.1:", when="@22.10: ^python@3.10")
    depends_on("py-typing-extensions@0.4.0.1:", when="@23.9: ^python@3.10")
    depends_on("py-uvloop@0.15.2:", when="@21.5-beta2:21,22.10:+uvloop")
```

# When would we go to “Version 1.0”?

---

Big things we've wanted for 1.0 are:

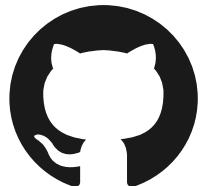
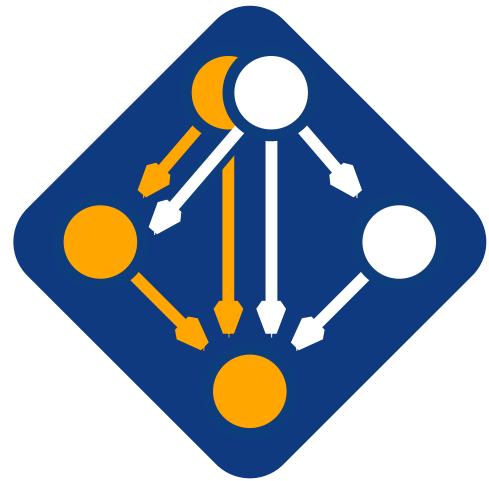
- New concretizer
  - production CI
  - production public build cache
  - Compilers as dependencies
  - Buildcache hardening
  - Stable package API
    - Enables separate package repository
- Done!**
- Aiming for November**
- Aiming for June**

We are getting very close!



# Join the Spack community!

- Join us and 3,200 others on [slack.spack.io](https://slack.spack.io)
- Contribute packages, documentation, and features on [GitHub](https://github.com/spack/spack)
- Continue the tutorial at [spack-tutorial.rtfd.io](https://spack-tutorial.rtfd.io)



★ Star us on GitHub!  
[github.com/spack/spack](https://github.com/spack/spack)



We hope to make distributing & using HPC software easy!

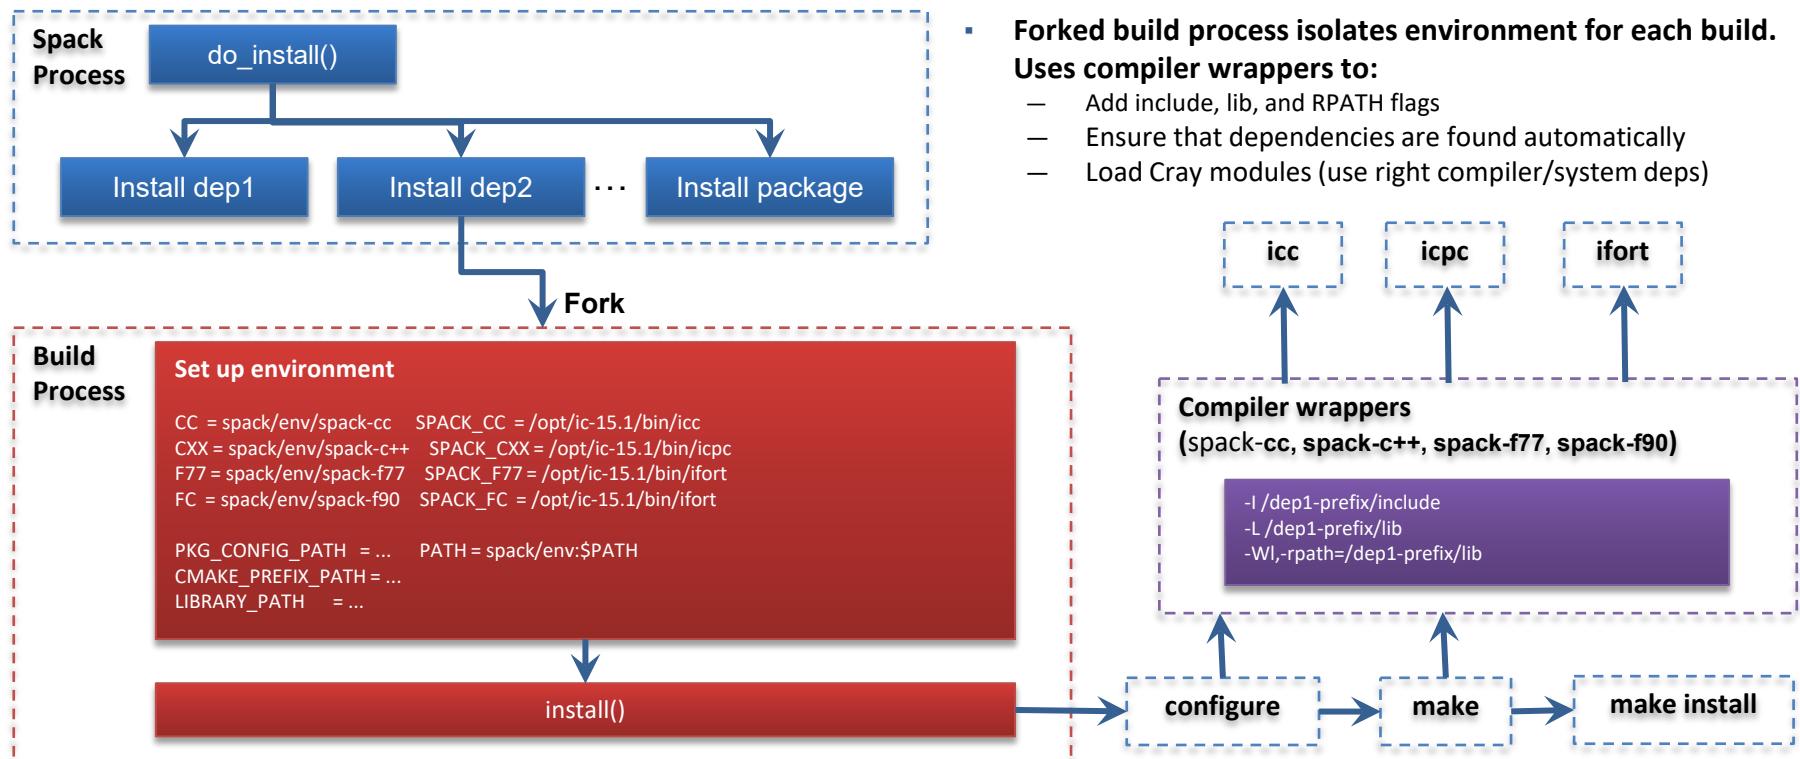


#### Disclaimer

This document was prepared as an account of work sponsored by an agency of the United States government. Neither the United States government nor Lawrence Livermore National Security, LLC, nor any of their employees makes any warranty, expressed or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States government or Lawrence Livermore National Security, LLC. The views and conclusions of authors expressed herein do not necessarily state or reflect those of the United States government or Lawrence Livermore National Security, LLC, and shall not be used for advertising or product endorsement purposes.



# An isolated compilation environment allows Spack to easily swap compilers



# OCI build caches

- Spack can now use arbitrary OCI registries as buildcaches

```
$ spack mirror add my_registry oci://user/image          # Dockerhub
$ spack mirror add my_registry oci://ghcr.io/haampie/spack-test    # GHCR
$ spack mirror set --push --oci-username ... --oci-password ... my_registry # set login
$ spack buildcache push my_registry [specs...]$ spack solve hdf5 ^cmake@3.0.1
```

- And you can optionally add a base image to get *runnable* images

```
$ spack buildcache push --base-image ubuntu:23.04 my_registry python
Pushed ... as [image]:python-3.11.2-65txfcpqbmpawclvtasuog4yzmxwaoia.spack

$ docker run --rm -it [image]:python-3.11.2-65txfcpqbmpawclvtasuog4yzmxwaoia.spack
```

# ASP has enabled us to improve Spack's error messages

```
$ spack solve hdf5 ^cmake@3.0.1
```

**==> Error: concretization failed for the following reasons:**

1. Cannot satisfy 'cmake@3.0.1'  
2. Cannot satisfy 'cmake@3.0.1'  
    required because hdf5 ^cmake@3.0.1 requested from CLI  
3. Cannot satisfy 'cmake@3.18:' and 'cmake@3.0.1'  
    required because hdf5 ^cmake@3.0.1 requested from CLI  
    required because hdf5 depends on cmake@3.18: when @1.13:  
    required because hdf5 ^cmake@3.0.1 requested from CLI  
4. Cannot satisfy 'cmake@3.12:' and 'cmake@3.0.1'  
    required because hdf5 depends on cmake@3.12:  
    required because hdf5 ^cmake@3.0.1 requested from CLI  
    required because hdf5 ^cmake@3.0.1 requested from CLI

# Spack's solver uses “generalized” conditions in ASP

- ASP rules: when the right side is true, the left is true

```
variant_value(Pkg, Variant, Value) :- variant_set(Pkg, Variant, Value), node(Pkg).
```

- Generalized conditions:
  - When condition requirements are True, the condition holds
  - When the condition holds, the imposed constraints are attributes of the spec

```
condition_holds(ID).  
:- attr(Name, A1), : condition_requirement(Name, A1);  
attr(Name, A1 , A2), : condition_requirement(Name, A1, A1).  
attr(Name, A1) :- condition_holds(ID), imposed_constraint(ID, Name, A1).  
attr(Name, A1, A2) :- condition_holds(ID), imposed_constraint(ID, Name, A1, A2).
```

- This is how all conditional logic in Spack is implemented

# We now build chains of causes and effects from the error message back to other sources of the problem

```
condition_cause(Effect, Cause) :-  
    condition_holds(Effect),  
    condition_holds(Cause),  
    attr(Name, A1),  
    condition_requirement(Effect, Name, A1),  
    imposed_constraint(Cause, Name, A1, A2).
```

Arity 1

```
condition_cause(Effect, Cause) :-  
    condition_holds(Effect),  
    condition_holds(Cause),  
    attr(Name, A1, A2),  
    condition_requirement(Effect, Name, A1, A2),  
    imposed_constraint(Cause, Name, A1, A2).
```

Arity 2

etc.

